

Algorithms for Dynamic Right-Sizing in Data Centers

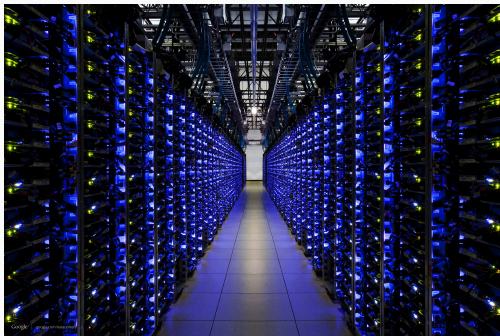
Kevin Kappelmann

August 03, 2017

Technical University of Munich

Dynamic Right-Sizing

Dynamically adjust the number of active servers and efficiently distribute the varying workloads



Model and Problem Formulation

Model Description

Input:

- $m \in \mathbb{N} \dots$ Number of homogeneous servers

Model Description

Input:

- $m \in \mathbb{N} \dots$ Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$ Switching costs of a server

Model Description

Input:

- $m \in \mathbb{N} \dots$ Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$ Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$ Convex operating cost function of a server

Model Description

Input:

- $m \in \mathbb{N} \dots$ Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$ Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$ Convex operating cost function of a server
- $T \in \mathbb{N} \dots$ Number of time slots

Model Description

Input:

- $m \in \mathbb{N} \dots$ Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$ Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$ Convex operating cost function of a server
- $T \in \mathbb{N} \dots$ Number of time slots
- $\lambda_1, \dots, \lambda_T \in [0, m] \dots$ Arrival rates (workloads)

Model Description

Input:

- $m \in \mathbb{N} \dots$ Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$ Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$ Convex operating cost function of a server
- $T \in \mathbb{N} \dots$ Number of time slots
- $\lambda_1, \dots, \lambda_T \in [0, m] \dots$ Arrival rates (workloads)

Notation:

- $x_1, \dots, x_T \in \{0, \dots, m\} \dots$ Numbers of active servers

Problem Statement

Operating costs for one time step t

$$c_{op}(x_t, \lambda_t) := \left\{ \right.$$

Problem Statement

Operating costs for one time step t

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t \end{cases} \quad // \text{too few servers}$$

Problem Statement

Operating costs for one time step t

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \end{cases}$$

Problem Statement

Operating costs for one time step t

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \\ x_t f(\lambda_t/x_t), & \text{otherwise} & // \text{even distribution} \end{cases}$$

Problem Statement

Operating costs for one time step t

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \\ x_t f(\lambda_t/x_t), & \text{otherwise} & // \text{even distribution} \end{cases}$$

Goal: Minimize total costs

Problem Statement

Operating costs for one time step t

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \\ x_t f(\lambda_t/x_t), & \text{otherwise} & // \text{even distribution} \end{cases}$$

Goal: Minimize total costs

$$\text{minimize} \quad \sum_{t=1}^T \underbrace{\left(c_{op}(x_t, \lambda_t) + \beta \max\{0, x_t - x_{t-1}\} \right)}_{c(x_{t-1}, x_t, \lambda_t)}$$

Optimal Offline Algorithm

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem

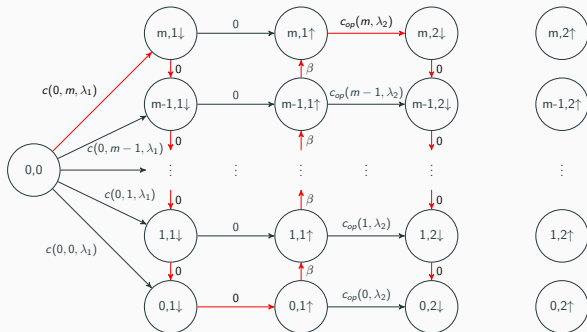


Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .


Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



Pseudo-Linear-Time Algorithm

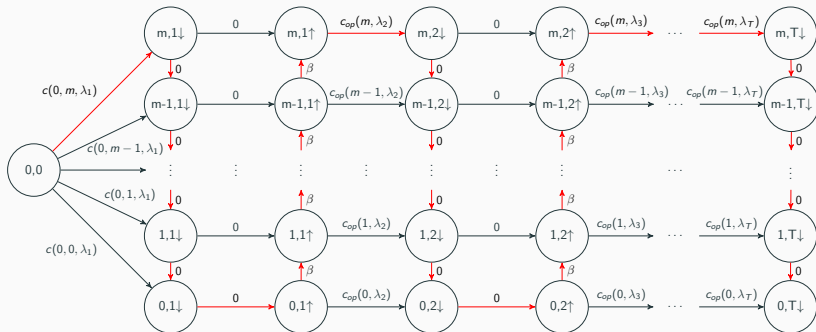
Fundamental idea: reduce problem to shortest path problem



Time complexity: $\Theta(Tm)$

Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



Time complexity: $\Theta(Tm)$

... only $\log_2(m)$ bits required to encode m .


Offline Approximation Algorithm

2-Optimal Linear-Time Algorithm

Use logarithmic steps to reduce number of nodes


2-Optimal Linear-Time Algorithm

Use logarithmic steps to reduce number of nodes



where $b := \lfloor \log_2(m) \rfloor$

2-Optimal Linear-Time Algorithm



Approximative Scheduling

Shortest Path in graph corresponds to 2-optimal schedule.

Approximative Scheduling

Shortest Path in graph corresponds to 2-optimal schedule.

Time complexity: $\Theta(T \log_2(m))$

Approximative Scheduling

Shortest Path in graph corresponds to 2-optimal schedule.

Time complexity: $\Theta(T \log_2(m))$

Can we generalize to allow for arbitrary precisions?

Approximative Scheduling

Shortest Path in graph corresponds to 2-optimal schedule.

Time complexity: $\Theta(T \log_2(m))$

Can we generalize to allow for arbitrary precisions? **Yes!**

Approximative Scheduling

Shortest Path in graph corresponds to 2-optimal schedule.

Time complexity: $\Theta(T \log_2(m))$

Can we generalize to allow for arbitrary precisions? Yes!

Approximation for arbitrary precision $\epsilon > 0$ with
time complexity $\Theta(T \log_{1+\epsilon}(m))$

Summary and Prospects

Summary

We reduced the scheduling problem to the shortest path problem of acyclic graphs.

Summary

We reduced the scheduling problem to the shortest path problem of acyclic graphs.

Optimal Offline Algorithm with pseudo-linear runtime
 $\Theta(Tm)$

Summary

We reduced the scheduling problem to the shortest path problem of acyclic graphs.

Optimal Offline Algorithm with pseudo-linear runtime
 $\Theta(Tm)$

$(1 + \epsilon)$ -Optimal Offline Algorithm with linear runtime
 $\Theta(T \log_{1+\epsilon}(m))$

Prospects

Can our approach be modified to ...

Can our approach be modified to ...

- deal with more than one homogeneous server collection?

Can our approach be modified to ...

- deal with more than one homogeneous server collection?
- deal with multiple sleep states?

Can our approach be modified to ...

- deal with more than one homogeneous server collection?
- deal with multiple sleep states?
- work as an online algorithm?

Can our approach be modified to ...

- deal with more than one homogeneous server collection?
- deal with multiple sleep states?
- work as an online algorithm?

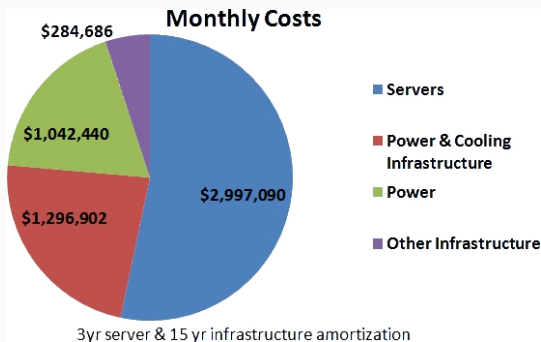
Open question: Is there a polynomial optimal algorithm or is it an NP-hard problem?

Thanks for your attention!

Any questions?

Data Centers' Costs

Fast-growing demand for data collection, processing, and storage

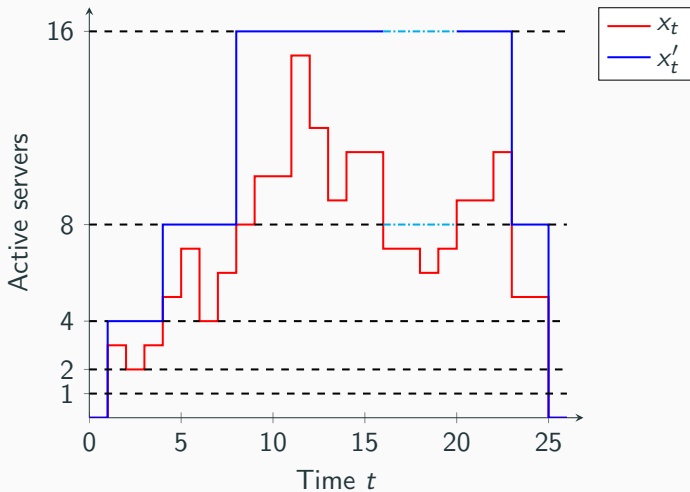


Proof Idea


Take a schedule and transform periods between two powers of 2

Proof Idea

Take a schedule and transform periods between two powers of 2



$(1 + \varepsilon)$ -Optimal Offline Algorithm



where $y := 1 + \varepsilon$, $b := \lfloor \log_y(m) \rfloor$

Image Sources I

- Data center: datacentervoice.com/wp-content/uploads/2015/12/data-center.jpg
- Data center costs: perspectives.mvdirona.com/2008/11/cost-of-power-in-large-scale-data-centers/