

# Algorithms for Dynamic Right-Sizing in Data Centers

---

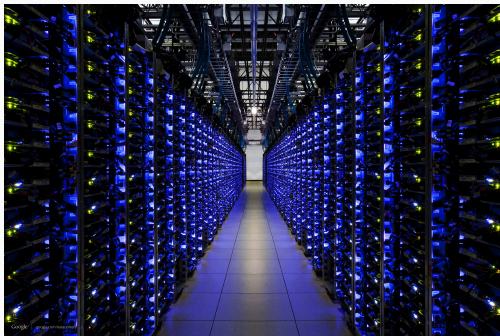
Kevin Kappelmann

August 03, 2017

Technical University of Munich

# Dynamic Right-Sizing

Dynamically adjust the number of active servers and efficiently distribute the varying workloads



# Model and Problem Formulation

---

# Model Description

Input:

- $m \in \mathbb{N} \dots$  Number of homogeneous servers

# Model Description

Input:

- $m \in \mathbb{N} \dots$  Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$  Switching costs of a server

# Model Description

Input:

- $m \in \mathbb{N} \dots$  Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$  Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$  Convex operating cost function of a server

# Model Description

Input:

- $m \in \mathbb{N} \dots$  Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$  Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$  Convex operating cost function of a server
- $T \in \mathbb{N} \dots$  Number of time slots

# Model Description

Input:

- $m \in \mathbb{N} \dots$  Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$  Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$  Convex operating cost function of a server
- $T \in \mathbb{N} \dots$  Number of time slots
- $\lambda_1, \dots, \lambda_T \in [0, m] \dots$  Arrival rates (workloads)



# Model Description

Input:

- $m \in \mathbb{N} \dots$  Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$  Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$  Convex operating cost function of a server
- $T \in \mathbb{N} \dots$  Number of time slots
- $\lambda_1, \dots, \lambda_T \in [0, m] \dots$  Arrival rates (workloads)

Conventions:

- $x_1, \dots, x_T \in \{0, \dots, m\} \dots$  Numbers of active servers

# Model Description

Input:

- $m \in \mathbb{N} \dots$  Number of homogeneous servers
- $\beta \in \mathbb{R}_{\geq 0} \dots$  Switching costs of a server
- $f : [0, 1] \rightarrow \mathbb{R} \dots$  Convex operating cost function of a server
- $T \in \mathbb{N} \dots$  Number of time slots
- $\lambda_1, \dots, \lambda_T \in [0, m] \dots$  Arrival rates (workloads)

Conventions:

- $x_1, \dots, x_T \in \{0, \dots, m\} \dots$  Numbers of active servers
- All servers are shut down before and after the scheduling process.

## Problem Statement

Operating costs for one time step  $t$

$$c_{op}(x_t, \lambda_t) := \left\{ \right.$$

## Problem Statement

Operating costs for one time step  $t$

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t \end{cases} \quad // \text{too few servers}$$

## Problem Statement

Operating costs for one time step  $t$

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \end{cases}$$

## Problem Statement

Operating costs for one time step  $t$

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \\ x_t f(\lambda_t/x_t), & \text{otherwise} & // \text{even distribution} \end{cases}$$

## Problem Statement

Operating costs for one time step  $t$

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \\ x_t f(\lambda_t/x_t), & \text{otherwise} & // \text{even distribution} \end{cases}$$

**Goal: Minimize total costs**

# Problem Statement

Operating costs for one time step  $t$

$$c_{op}(x_t, \lambda_t) := \begin{cases} \infty, & \text{if } \lambda_t > x_t & // \text{too few servers} \\ 0, & \text{if } x_t = \lambda_t = 0 & // \text{no active servers} \\ x_t f(\lambda_t/x_t), & \text{otherwise} & // \text{even distribution} \end{cases}$$

**Goal: Minimize total costs**

$$\text{minimize} \quad \sum_{t=1}^T \underbrace{\left( c_{op}(x_t, \lambda_t) + \beta \max\{0, x_t - x_{t-1}\} \right)}_{c(x_{t-1}, x_t, \lambda_t)}$$



# Optimal Offline Algorithm

---

# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem

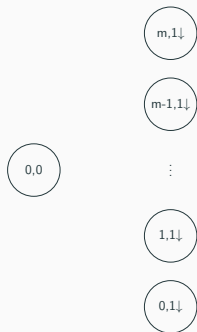
# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



# Pseudo-Linear-Time Algorithm


Fundamental idea: reduce problem to shortest path problem





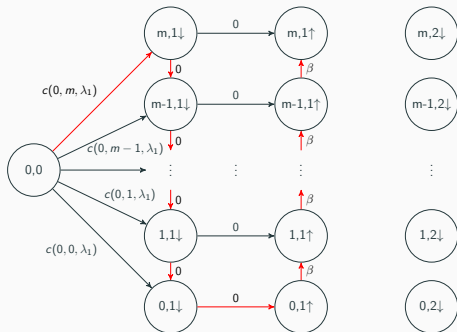
# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



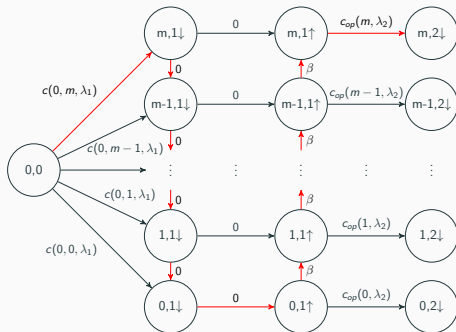
# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem




# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



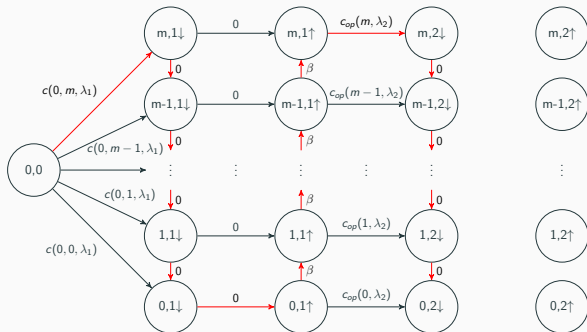
# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



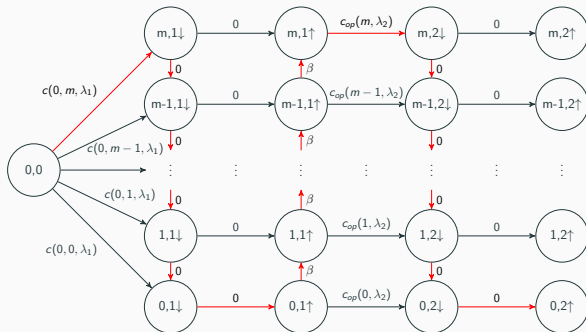
# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



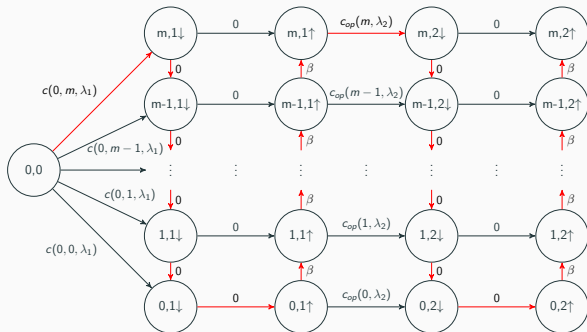
# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



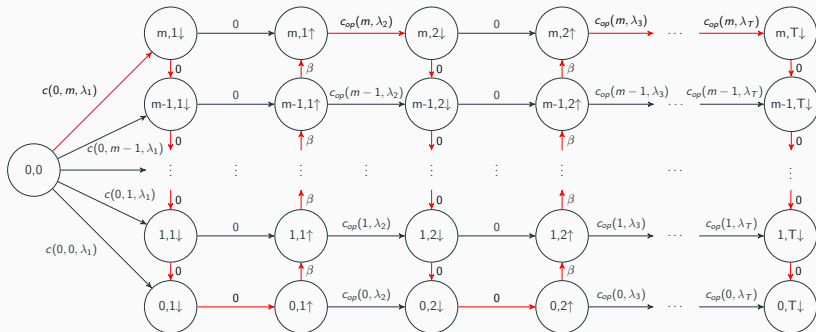
# Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



# Pseudo-Linear-Time Algorithm


Fundamental idea: reduce problem to shortest path problem





## Pseudo-Linear-Time Algorithm


Fundamental idea: reduce problem to shortest path problem



Time complexity:  $\Theta(Tm)$

## Pseudo-Linear-Time Algorithm

Fundamental idea: reduce problem to shortest path problem



Time complexity:  $\Theta(Tm)$

... only  $\log_2(m)$  bits required to encode  $m$ .

# Offline Approximation Algorithm


---

## 2-Optimal Linear-Time Algorithm

Use logarithmic steps to reduce number of nodes

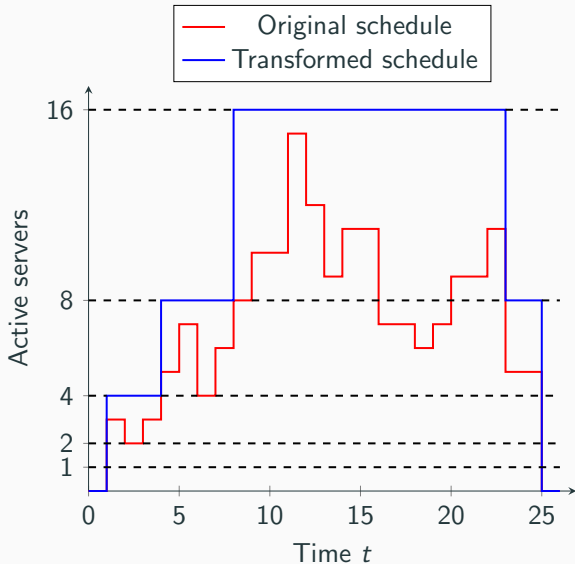
## 2-Optimal Linear-Time Algorithm

Use logarithmic steps to reduce number of nodes



where  $b := \lfloor \log_2(m) \rfloor$

## 2-Optimal Linear-Time Algorithm



## Approximative Scheduling

A shortest path in the graph corresponds to a 2-optimal schedule.

## Approximative Scheduling

A shortest path in the graph corresponds to a 2-optimal schedule.

Time complexity:  $\Theta(T \log_2(m))$



## Approximative Scheduling

A shortest path in the graph corresponds to a 2-optimal schedule.

Time complexity:  $\Theta(T \log_2(m))$

Can we generalize to allow for arbitrary precisions?

## Approximative Scheduling

A shortest path in the graph corresponds to a 2-optimal schedule.

Time complexity:  $\Theta(T \log_2(m))$

Can we generalize to allow for arbitrary precisions? **Yes!**

## Approximative Scheduling

A shortest path in the graph corresponds to a 2-optimal schedule.

Time complexity:  $\Theta(T \log_2(m))$

Can we generalize to allow for arbitrary precisions? Yes!

Approximation for arbitrary precision  $\epsilon > 0$  with  
time complexity  $\Theta(T \log_{1+\epsilon}(m))$

# Summary and Prospects

---

## Summary

We reduced the scheduling problem to the shortest path problem of acyclic graphs.

## Summary

We reduced the scheduling problem to the shortest path problem of acyclic graphs.

Optimal Offline Algorithm with pseudo-linear runtime  
 $\Theta(Tm)$

## Summary

We reduced the scheduling problem to the shortest path problem of acyclic graphs.

Optimal Offline Algorithm with pseudo-linear runtime  
 $\Theta(Tm)$

$(1 + \epsilon)$ -Optimal Offline Algorithm with linear runtime  
 $\Theta(T \log_{1+\epsilon}(m))$

# Prospects

Can our approach be modified to ...



Can our approach be modified to ...

- deal with more than one homogeneous server collection?

Can our approach be modified to ...

- deal with more than one homogeneous server collection?
- deal with multiple sleep states?

Can our approach be modified to ...

- deal with more than one homogeneous server collection?
- deal with multiple sleep states?
- work as an online algorithm?

Can our approach be modified to ...

- deal with more than one homogeneous server collection?
- deal with multiple sleep states?
- work as an online algorithm?

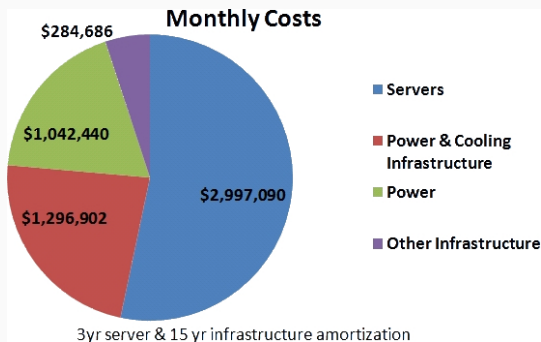
Open question: Is there a polynomial-time optimal algorithm or is it an NP-hard problem?

**Thanks for your attention!**

**Any questions?**

# Data Centers' Costs

Fast-growing demand for data collection, processing, and storage

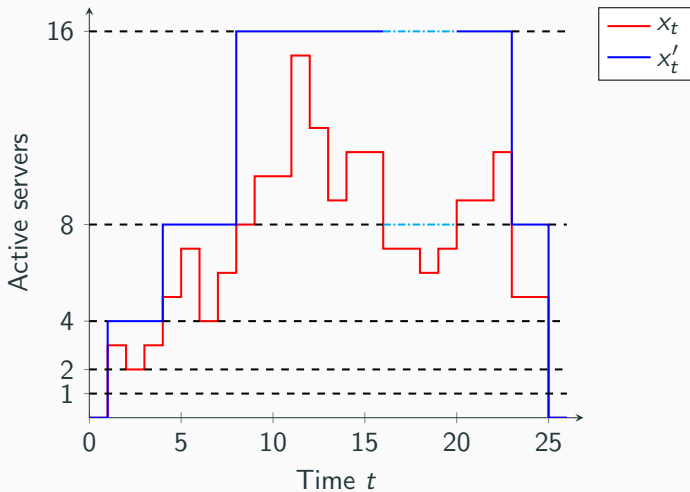


## Proof Idea

Take a schedule and transform periods between two powers of 2

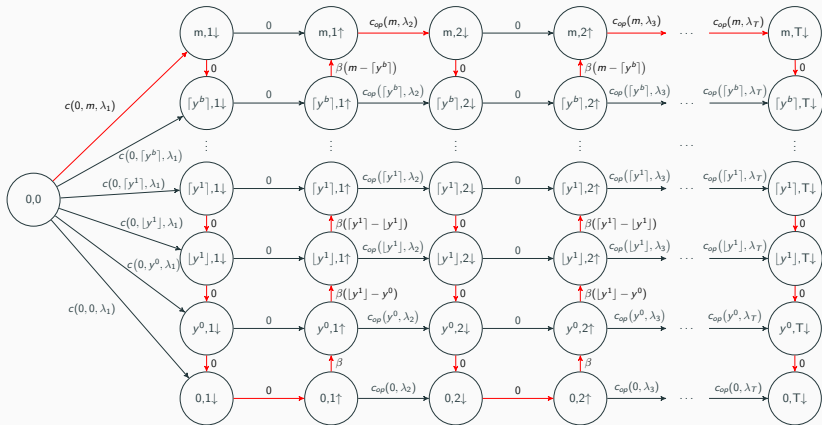
# Proof Idea

Take a schedule and transform periods between two powers of 2





# $(1 + \varepsilon)$ -Optimal Offline Algorithm



where  $y := 1 + \varepsilon$ ,  $b := \lfloor \log_y(m) \rfloor$

## Image Sources I

- Data center: [datacentervoice.com/wp-content/uploads/2015/12/data-center.jpg](http://datacentervoice.com/wp-content/uploads/2015/12/data-center.jpg)
- Data center costs: [perspectives.mvdirona.com/2008/11/cost-of-power-in-large-scale-data-centers/](http://perspectives.mvdirona.com/2008/11/cost-of-power-in-large-scale-data-centers/)