

1 Introduction

1.1 Input and conventions

TODO: rewrite in nice paragraph

Input:

- $m \in \mathbb{N}$: Number of homogeneous servers
- $T \in \mathbb{N}$: Number of time steps
- $\beta \in \mathbb{R}_{\geq 0}$: Power up costs
- $\lambda_0, \dots, \lambda_T \in [0, m]$: Arrival rates

Notations:

- Let $\lambda_{i,t}$ be the assigned arrival rate at time t for server i
- Let x_t be the number of active servers at time t
- Let $\mathcal{X} := (x_0, \dots, x_T)$ be the sequence of active servers
- If $A = (a_0, \dots, a_n)$ is a tuple with $n + 1$ entries, we write $A(i)$ for the i -th component of A with $0 \leq i \leq n$

Requirements:

- Convex cost function f
- Power down costs are w.l.o.g. equal to 0
- $\forall t \in \{0, \dots, T\} : \sum_{i=1}^m \lambda_{i,t} = \lambda_t$
- $\lambda_0 = \lambda_T = 0$
- $\mathcal{X}(0) = \mathcal{X}(T) = 0$, i.e. all servers are powered down at $t = 0$ and $t = T$

1.2 Preliminaries

Lemma 1.1. *Given a convex cost function f , x active servers and an arrival rate λ , the best method is to assign each server a load of λ/x .*

Proof. $\forall x \in \mathbb{N}, \mu_i \in [0, 1] : \sum_{i=1}^x \mu_i = 1 :$

$$\begin{aligned} f\left(\frac{\lambda}{x}\right) &= f\left(\sum_{i=1}^x \frac{\mu_i * \lambda}{x}\right) \stackrel{\text{Jensen's inequality}}{\leq} \sum_{i=1}^x \frac{1}{x} f(\mu_i * \lambda) \\ &\Leftrightarrow x * f\left(\frac{\lambda}{x}\right) \leq \sum_{i=1}^x f(\mu_i * \lambda) \end{aligned}$$

□

Lemma 1.1 tells us, that having x_t active servers, it is always the best method to equally share λ_t on all x_t active servers. This allows us to uniquely identify an optimal schedule by the sequence of numbers of active servers \mathcal{X} .

The costs of a schedule are then given by:

$$costs(\mathcal{X}) := \sum_{t=0}^T \underbrace{\beta * \max\{0, x_t - x_{t-1}\}}_{\text{power up costs}} + x_t * f(\lambda_t/t)$$

We call a schedule \mathcal{X} **feasible** if $\forall t \in \{0, \dots, T\} : x_t \geq \lambda_t$. In particular, every optimal schedule is feasible.

2 Optimal scheduling for m homogeneous servers

TODO: introduction text

2.1 Graph for an optimal schedule

We construct a directed acyclic graph as follows:

$\forall t \in [T - 1]$ and $i, j \in \{0, \dots, m\}$ we add vertices (t, i) modelling the number of active servers at time t . Furthermore, we add vertices $(0, 0)$ and $(T, 0)$ for our initial and final state respectively. In order to warrant that there are at least $\lceil \lambda_t \rceil$ active servers $\forall t \in [T - 1]$, we define an auxiliary function which calculates the costs for handling an arrival rate λ with x active servers:

$$c(x, \lambda) := \begin{cases} 0 & \text{if } x = 0 \\ x * f(\lambda/x), & \text{if } x \neq 0 \wedge \lambda \leq x \\ \infty, & \text{otherwise} \end{cases} \quad (1)$$

Then, $\forall t \in [T - 2]$ and $i, j \in \{0, \dots, m\}$ we add edges from (t, i) to $(t + 1, j)$ with weight

$$d(i, j, \lambda_{t+1}) := \beta * \max\{0, j - i\} + c(j, \lambda_{t+1}) \quad (2)$$

Finally, for $0 \leq i \leq m$ we add edges from $(0, 0)$ to $(1, i)$ with weight $d(0, i, \lambda_1)$ and from $(T - 1, i)$ to $(T, 0)$ with weight $d(i, 0, \lambda_T) = 0$.

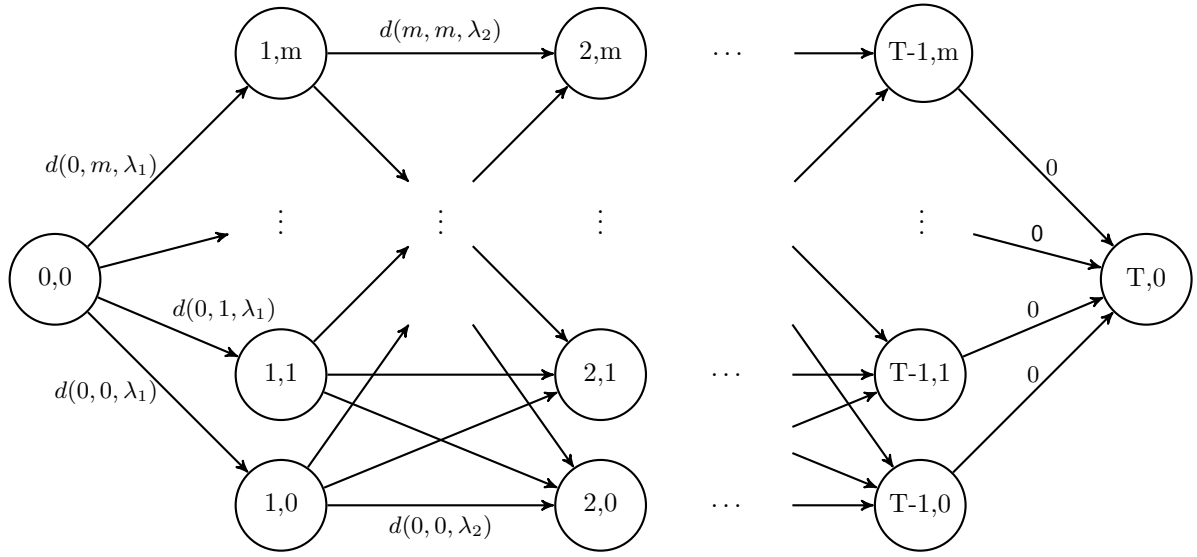


Figure 1: Graph for optimal schedule algorithm.

Note: All edges from (t, i) to $(t + 1, j)$ have weight $d(i, j, \lambda_{t+1})$

Proposition 2.1. Any given optimal schedule \mathcal{X} corresponds to a shortest path P from $(0, 0)$ to $(T, 0)$ with $\text{costs}(\mathcal{X}) = \text{costs}(P)$ and vice versa.

Proof.

“ \Rightarrow ”: We construct a feasible path in our graph from \mathcal{X} as follows:

$\forall t \in [T]$ set $e_t := ((t - 1, \mathcal{X}(t - 1)), (t, \mathcal{X}(t)))$. Then set $P := (e_1, \dots, e_T)$.

As each edge e_t in our graph has weight $d(\mathcal{X}(t - 1), \mathcal{X}(t), \lambda_t)$ and hence corresponds to the costs of switching from $\mathcal{X}(t - 1)$ to $\mathcal{X}(t)$ servers and processing λ_t with $\mathcal{X}(t)$ active servers, it directly follows that P is a shortest path of the graph with $\text{costs}(P) = \text{costs}(\mathcal{X})$.

“ \Leftarrow ”: Let $P = ((0, 0) = v_0, \dots, v_T = (T, 0))$ with $v_t \in \{(t, i) \mid 0 \leq i \leq m\}$ be a shortest path of the graph.

We can construct an optimal schedule from P by setting $\mathcal{X} := (v_0(1), \dots, v_T(1))$

By definition (1) it is guaranteed that P only traverses edges such that there are enough active servers $\forall t \in [T]$. Therefore, the created schedule is feasible. It's optimality directly follows from the definition of the edges' weights and so does the equality $\text{costs}(\mathcal{X}) = \text{costs}(P)$.

□

2.2 A minimum cost algorithm

Algorithm 1 Calculate costs for m homogeneous servers

Require: Convex cost function f , $\lambda_0 = \lambda_T = 0$, $\forall t \in [T-1] : \lambda_t \in [0, m]$

```

1: function SCHEDULE( $m, T, \beta, \lambda_1, \dots, \lambda_{T-1}$ )
2:   if  $T < 2$  then return
3:   let  $p[2 \dots T-1, m]$  and  $M[1 \dots T-1, m]$  be new arrays
4:   for  $j \leftarrow 0$  to  $m$  do
5:      $M[1, j] \leftarrow d(0, j, \lambda_1)$ 
6:   for  $t \leftarrow 1$  to  $T-2$  do
7:     for  $j \leftarrow 0$  to  $m$  do
8:        $opt \leftarrow \infty$ 
9:       for  $i \leftarrow 0$  to  $m$  do
10:         $M[t+1, j] \leftarrow M[t, i] + d(i, j, \lambda_{t+1})$ 
11:        if  $M[t+1, j] < opt$  then
12:           $opt \leftarrow M[t+1, j]$ 
13:           $p[t+1, j] \leftarrow i$ 
14:         $M[t+1, j] \leftarrow opt$ 
15:   return  $p$  and  $M$ 

```

Algorithm 2 Extract schedule for n homogeneous servers

```

1: function EXTRACT( $m, p, M, T$ )
2:   let  $x[0 \dots T]$  be a new array
3:    $x[0] \leftarrow x[T] \leftarrow 0$ 
4:   if  $T < 2$  then return  $x$   $\triangleright$  Trivial solution
5:    $x[T-1] \leftarrow \arg \min_{0 \leq i \leq m} \{M[T-1, i]\}$ 
6:   for  $t \leftarrow T-2$  to  $1$  do
7:      $x[t] \leftarrow p[t+1, x[t+1]]$ 
8:   return  $x$ 

```

2.2.1 Runtime analysis

Schedule: Loop 5,8 and 10 run $m+1$ times, loop 7 runs $T-2$ times

Extract: Loop 5 runs $T-2$ times, argmin 4 takes time $m+1$.

For $T, n \rightarrow \infty$ it holds:

$$\mathcal{O}(m+1 + (T-2) * (m+1)^2 + T-2 + m+1) = \mathcal{O}(2 * m + T + (T-2) * (m+1)^2) = \mathcal{O}(T * m^2) \quad (3)$$

As we need $\log_2(m)$ bits to encode m , the algorithm is exponential in the number of servers.

2.2.2 A memory optimized algorithm

TODO

3 A 2-approximation algorithm for monotonically increasing convex f

We consider a modification of the problem discussed in chapter 2. Assuming that f is convex and monotonically increasing, we can modify our algorithm to obtain a polynomial time 2-approximation algorithm.

3.1 Graph for a 2-optimal schedule

We modify our graph from chapter 2.1 to the reduce the number of vertices. For this, we stop adding m vertices for each timestep but using vertices that approximate the number of active servers. First, let $b := \lceil \log_2(m) \rceil$. We add vertices $(t, 0)$ and $(t, 2^i)$, $\forall t \in [T-1], 0 \leq i \leq b$. All edges and weights are added analogous to chapter 2.1.




Figure 2: Graph for 2-approximation algorithm

Proposition 3.1.

1. Any given optimal schedule \mathcal{X} can be transformed to a 2-optimal schedule \mathcal{X}' which corresponds to a path P from $(0, 0)$ to $(T, 0)$ with $\text{costs}(\mathcal{X}') = \text{costs}(P)$.
2. Any shortest path P from $(0, 0)$ to $(T, 0)$ corresponds to a 2-optimal schedule \mathcal{X}' with $\text{costs}(P) = \text{costs}(\mathcal{X}')$.

Proof.

1. Assume we have an optimal schedule identified by $\mathcal{X} = (x_0, \dots, x_T)$. For $0 \leq t \leq T$ we set:

$$x'_t := \begin{cases} 2^{\lceil \log_2(x_t) \rceil}, & \text{if } x_t > 0 \\ 0, & \text{otherwise} \end{cases} \quad (4)$$

Now let $\mathcal{X}' := (x'_0, \dots, x'_T)$ be the modified sequence of active servers. Notice that $x_t \leq x'_t \leq 2 * x_t$ holds as x'_t is the smallest power of two larger or equal to x_t and therefore \mathcal{X}' is

feasible. We can construct a feasible path in our graph from \mathcal{X}' as follows:

$\forall t \in [T]$ set $e_t := \left((t-1, \mathcal{X}'(t-1)), (t, \mathcal{X}'(t)) \right)$. Then set $P := (e_1, \dots, e_T)$. By the definition of the edges' weights it follows $\text{costs}(\mathcal{X}') = \text{costs}(P)$.

Next, let's consider the difference of costs between \mathcal{X}' and \mathcal{X} at every time step t to $t+1$ for $0 \leq t \leq T-1$:

(a) $x_{t+1} > 0$: The costs incurred by \mathcal{X}' are given by

$$\begin{aligned}
d(\mathcal{X}'(t), \mathcal{X}'(t+1), \lambda_{t+1}) &= \beta * \max\{0, x'_{t+1} - x'_t\} + c(x'_{t+1}, \lambda_{t+1}) \\
(\mathcal{X}' \text{ is feasible}) &= \beta * \max\{0, x'_{t+1} - x'_t\} + x'_{t+1} * f(\lambda_{t+1}/x'_{t+1}) \\
(4) &\leq \beta * \max\{0, x'_{t+1} - x'_t\} + 2 * x_{t+1} * f(\lambda_{t+1}/x'_{t+1}) \\
(\text{f monotonically increasing}) &\leq \beta * \max\{0, x'_{t+1} - x'_t\} + 2 * x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \quad (5)
\end{aligned}$$

and the costs of \mathcal{X} by

$$\begin{aligned}
d(\mathcal{X}(t), \mathcal{X}(t+1), \lambda_{t+1}) &= \beta * \max\{0, x_{t+1} - x_t\} + c(x_{t+1}, \lambda_{t+1}) \\
(\mathcal{X} \text{ is feasible}) &= \beta * \max\{0, x_{t+1} - x_t\} + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \quad (6)
\end{aligned}$$

Therefore, we can estimate the difference:

$$\begin{aligned}
&d(\mathcal{X}'(t), \mathcal{X}'(t+1), \lambda_{t+1}) - d(\mathcal{X}(t), \mathcal{X}(t+1), \lambda_{t+1}) \\
(5),(6) &\leq \beta * \max\{0, x'_{t+1} - x'_t\} - \beta * \max\{0, x_{t+1} - x_t\} + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
&= \beta * (\max\{0, x'_{t+1} - x'_t\} - \max\{0, x_{t+1} - x_t\}) + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \quad (7)
\end{aligned}$$

At this point, we need another case distinction.

(i) $x_{t+1} \leq x_t$: From (4) it follows that $x'_{t+1} \leq x'_t$. We continue to simplify (7):

$$\begin{aligned}
&\beta * (\max\{0, x'_{t+1} - x'_t\} - \max\{0, x_{t+1} - x_t\}) + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
&= \beta * (0 - 0) + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
&= x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
&= d(\mathcal{X}(t), \mathcal{X}(t+1), \lambda_{t+1})
\end{aligned}$$

$$\Rightarrow \text{costs}(\mathcal{X}') = 2 * \text{costs}(\mathcal{X})$$

(ii) $x_{t+1} > x_t$: We simplify (7) and obtain

$$\begin{aligned}
&\beta * (\max\{0, x'_{t+1} - x'_t\} - \max\{0, x_{t+1} - x_t\}) + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
&= \beta * (x'_{t+1} - x'_t - x_{t+1} + x_t) + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
(4) &\leq \beta * (2 * x_{t+1} - x_t - x_{t+1} + x_t) + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
&= \beta * x_{t+1} + x_{t+1} * f(\lambda_{t+1}/x_{t+1}) \\
&= d(0, \mathcal{X}(t+1), \lambda_{t+1}) \\
&\neq d(\mathcal{X}(t), \mathcal{X}(t+1), \lambda_{t+1})
\end{aligned}$$

Here we fail.

(b) $x_{t+1} = 0$: Then by (4) we conclude that $x_{t+1} = x'_{t+1} = 0$ and therefore:

$$\begin{aligned}
&d(\mathcal{X}'(t), \mathcal{X}'(t+1), \lambda_{t+1}) - d(\mathcal{X}(t), \mathcal{X}(t+1), \lambda_{t+1}) \\
&= \beta * \max\{0, 0 - x'_t\} + c(0, \lambda_{t+1}) - \beta * \max\{0, 0 - x_t\} - c(0, \lambda_{t+1}) \\
&= 0
\end{aligned}$$

$$\Rightarrow \text{costs}(\mathcal{X}') = \text{costs}(\mathcal{X})$$

2. From (1) we obtain that we can construct a 2-optimal path P' from any optimal schedule. Now, let P be a shortest path. We have $costs(P) \leq costs(P')$ and as every feasible path in P corresponds to an feasible schedule \mathcal{X} with $costs(P) = costs(\mathcal{X})$, P must also be at least 2-optimal.

□