

Algorithms for Dynamic Right-Sizing in Data Centers

Kevin Kappelmann

*Chair for Theoretical Computer Science,
Technical University of Munich*

May 9, 2017

Contents

1	Introduction	1
1.1	Model description	1
1.2	Problem statement	1
2	Preliminaries	2
3	Optimal offline scheduling	5
3.1	Graph for an optimal schedule	5
3.2	A pseudo-polynomial minimum cost algorithm	7
3.2.1	Runtime analysis	7
3.2.2	A memory optimized algorithm	7
4	A polynomial 4-approximation algorithm for monotonically increasing convex f	8
4.1	Graph for a 4-optimal schedule	8

1 Introduction

TODO: Hardware prices vs. energy costs in data centres and purpose of this paper (offline algorithm, approximation algorithm,...).

1.1 Model description

We want to address the issue of above-mentioned ever-growing energy consumption by examining a scheduling problem that commonly arises in data centres. More specifically, we consider a model consisting of a fixed amount of homogeneous servers denoted by $m \in \mathbb{N}$ and a fixed amount of time slots denoted by $T \in \mathbb{N}$. In turn, each server possesses two power states, i.e. each server is either powered on (*active state*) or powered off (*sleep state*).

Need to describe what a time slot means?

Better name than sleep state?

For any time slot $t \in [T]$, we have a *mean arrival rate* denoted by λ_t , i.e. the amount of expected load to process in time slot t . We expect the arrival rates to be normalised such that each server $i \in [m]$ can handle a load between zero and one in any time slot. We denote the assigned load for server i in time slot t by $\lambda_{i,t} \in [0, 1]$. Consequently, for any time slot t , we expect an arrival rate between zero and m , i.e. $\lambda_t \in [0, m]$; otherwise, the servers would not be able to process the given load in time.

The incurring energy costs of a single machine can be described by the sum of the machine's (*power state*) *switching costs*, specified by $\beta \in \mathbb{R}_{\geq 0}$, as well as its *operating costs*, specified by $f : [0, 1] \rightarrow \mathbb{R}$. We assume that a sleeping server does not generate any energy costs. Note that $f(0)$ describes the costs generated by an idle server, not a sleeping one; in particular, $f(0)$ may be non-zero. Further, we assume convexity for f . This may seem like a notable restriction at first, but it indeed captures the behaviour of most modern server models. Since we are dealing with homogeneous servers, β and f are the same for all machines.

We want to stress that f may not exclusively consider energy costs. For example, f may also allow for costs incurred by delays, such as lost revenue caused by users waiting for their responses. Similarly, β may also allow for delay costs, wear and tear costs or the like. [1]

For convenience, we assume all machines sleeping at time $t = 0$ and force all machines to sleep after the scheduling process, i.e. at times $t > T$. This justifies the consolidation of power up and power down costs into β because it allows us to model both costs as being incurred when powering up a server; that is, a model with power up costs β_{\uparrow} and power down costs β_{\downarrow} can be simply transferred to our model by setting $\beta := \beta_{\uparrow} + \beta_{\downarrow}$.

1.2 Problem statement

Using above definitions, we can define the input of our model by setting $\mathcal{I} := (m, T, \Lambda, \beta, f)$ where $\Lambda = (\lambda_1, \dots, \lambda_T)$ is the sequence of arrival rates. Naturally, given an input \mathcal{I} , we want to schedule our servers in such a way that we minimise the sum of incurred costs while warranting that we are processing the given loads in time.

For this, consider for each server $i \in [m]$ the sequence of its states S_i and the sequence of

its assigned loads L_i ; that is

$$S_i := (s_{i,1}, \dots, s_{i,T}) \in \{0, 1\}^T$$

$$L_i := (\lambda_{i,1}, \dots, \lambda_{i,T}) \in [0, 1]^T$$

where $s_{i,t} \in \{0, 1\}$ denotes whether server i at time t is sleeping (0) or active (1). Recall that we assume all machines sleeping at times $t \notin [T]$; consequently, for $t \notin [T]$ and $i \in [m]$, we have $s_{i,t} = 0$.

Using these sequences S_i and L_i , we can now define the sequence of all state changes and the sequence of all assigned loads:

good sentence?
 $\in (\{0, 1\}^m)^T$

$$\mathcal{S} := (S_1, \dots, S_m)$$

$$\mathcal{L} := (L_1, \dots, L_m)$$

We will subsequently call a pair $\Sigma := (\mathcal{S}, \mathcal{L})$ a *schedule*. Finally, we are ready to define our problem statement.

Given an input \mathcal{I} , our goal is to find a schedule Σ that satisfies the following optimisation:

Definition of $c(\Sigma)$
too hidden?

$$\text{minimise} \quad c(\Sigma) := \underbrace{\sum_{t=1}^T \sum_{i=1}^m (f(\lambda_{i,t}) * s_{i,t})}_{\text{operating costs}} + \beta * \underbrace{\sum_{t=1}^T \sum_{i=1}^m \min\{0, s_{i,t} - s_{i,t-1}\}}_{\text{switching costs}} \quad (1)$$

$$\text{subject to} \quad \sum_{i=1}^m (\lambda_{i,t} * s_{i,t}) = \lambda_t, \quad \forall t \in [T] \quad (2)$$

We call a schedule *feasible* if it satisfies (2) and *optimal* if it satisfies (1) and (2).

2 Preliminaries

In this section, we conduct the preparatory work needed for our algorithms and proofs.

This sounds strange... What should be written here?

Proposition 2.1. *Given a problem instance \mathcal{I} and a feasible schedule Σ , there exists a feasible schedule Σ' such that*

(i) $c(\Sigma') \leq c(\Sigma)$ and

(ii) Σ' never powers on and shuts down servers at the same time slot, i.e. Σ' satisfies the following formula:

$$\forall t \in [T] \left[(\forall i \in [m] (s_{i,t} - s_{i,t-1} \geq 0)) \vee (\forall i \in [m] (s_{i,t} - s_{i,t-1} \leq 0)) \right] \quad (3)$$

Proof. Let $\Sigma = (\mathcal{S}, \mathcal{L})$ be a feasible schedule for \mathcal{I} . We give a procedure that repeatedly modifies Σ such that it satisfies (3) and reduces or retains its costs.

Let $t \in [T]$ be the first time slot falsifying (3). If there does not exist such a t , we are done. Otherwise, we can obtain $i, j \in [m]$ such that $s_{i,t} - s_{i,t-1} = 1$ and $s_{j,t} - s_{j,t-1} = -1$. WLOG we may assume $i < j$.

First, since all servers are sleeping at time $t = 0$, we have

$$s_{k,1} - s_{k,0} = s_{k,1} - 0 = s_{k,1} \geq 0, \quad \forall k \in [m]$$

Thus, we may assume $t > 1$. Now consider the state sequences S_i and S_j :

$$\begin{aligned} S_i &= (s_{i,1}, \dots, s_{i,t-1} = 0, s_{i,t} = 1, \dots, s_{i,T}) \\ S_j &= (s_{j,1}, \dots, s_{j,t-1} = 1, s_{j,t} = 0, \dots, s_{j,T}) \end{aligned}$$

We modify S_i and S_j by swapping their states for time slots $\geq t$, that is we set

$$\begin{aligned} S'_i &:= (s_{i,1}, \dots, s_{i,t-1} = 0, s_{j,t} = 0, \dots, s_{j,T}) \\ S'_j &:= (s_{j,1}, \dots, s_{j,t-1} = 1, s_{i,t} = 1, \dots, s_{i,T}) \end{aligned}$$

Similarly, we need to swap the assigned loads:

$$\begin{aligned} L'_i &:= (\lambda_{i,1}, \dots, \lambda_{i,t-1}, \lambda_{j,t}, \dots, \lambda_{j,T}) \\ L'_j &:= (\lambda_{j,1}, \dots, \lambda_{j,t-1}, \lambda_{i,t}, \dots, \lambda_{i,T}) \end{aligned}$$

Finally, we construct a new schedule $\Sigma' := (\mathcal{S}', \mathcal{L}')$ given by

$$\begin{aligned} \mathcal{S}' &:= (S_1, \dots, S'_i, \dots, S'_j, \dots, S_T) \\ \mathcal{L}' &:= (L_1, \dots, L'_i, \dots, L'_j, \dots, L_T) \end{aligned}$$

We now want to verify that Σ' is a feasible schedule, that is Σ' satisfies (2). For time slots $< t$ the schedules Σ' and Σ still coincide. For time slots $\geq t$ we only changed the order of summation (2). Thus, Σ' is feasible.

Further, Σ' and Σ coincide in their operating costs; however, their switching costs differ in that there are no switching costs at time slot t for server i using Σ' . As we assume $\beta \geq 0$, we conclude $c(\Sigma') \leq c(\Sigma)$.

Moreover, we decreased the amount of bad spots at time slot t concerning (3). Hence, by repeating described process on Σ' , we obtain a terminating procedure that returns a schedule satisfying the conditions. \square

bad spots? better description?

Corollary 2.2. *Given a problem instance \mathcal{I} , there exists an optimal schedule Σ satisfying (3).*

Proof. Let Σ be an optimal schedule for \mathcal{I} . Applying proposition 2.1 to Σ yields the result. \square

Next, we want to consider the sequence of active servers. For this, let \mathcal{X} denote the sequence of sums of active servers at each time slot t , i.e.

$$\mathcal{X} := (x_1 = \sum_{i=1}^m s_{i,1}, \dots, x_T = \sum_{i=1}^m s_{i,T}) \in \{0, \dots, m\}^T$$

Recall that we assume all machines sleeping at times $t \notin [T]$; consequently, for $t \notin [T]$ we have $x_t = 0$.

The next proposition poses the cornerstone of our subsequent works.

Proposition 2.3 (Equal load sharing). *Given $x_t \in \mathbb{N}$ active servers at time slot t , an arrival rate $\lambda_t \in [0, x_t]$, and a convex cost function f , a most cost-efficient and feasible scheduling strategy is to assign each active server a load of λ_t/x_t .*

This is sounds awkwardly formulated, doesn't it?

Proof. Let Σ be an arbitrary, feasible schedule using x_t servers at time slot t , and let A be its set of active servers at time slot t , that is $A := \{i \in [m] \mid s_{i,t} = 1\}$. Consider the operating costs of Σ given by

$$\sum_{i=1}^m (f(\lambda_{i,t}) * s_{i,t}) = \sum_{i \in A} f(\lambda_{i,t})$$

Since Σ is feasible (see constraint (2)), we have

$$\sum_{i \in A} \lambda_{i,t} = \lambda_t$$

Thus, we can obtain weights $\mu_1, \dots, \mu_{x_t} \in [0, 1]$ that relate $\lambda_{i,t}$ and λ_t for $i \in A$ such that

$$\sum_{i=1}^{x_t} \mu_i = 1 \quad \text{and} \quad \sum_{i \in A} f(\lambda_{i,t}) = \sum_{i=1}^{x_t} f(\mu_i \lambda_t)$$

In particular, we have $\sum_{i=1}^{x_t} \mu_i \lambda_t = \lambda_t$. Using these weights, we now consider the operating costs of a schedule Σ^* that equally distributes λ_t to its x_t active servers:

$$\sum_{i=1}^{x_t} f\left(\frac{\lambda_t}{x_t}\right) = x_t * f\left(\frac{\lambda_t}{x_t}\right) = x_t * f\left(\sum_{i=1}^{x_t} \frac{\mu_i \lambda_t}{x_t}\right)$$

Using Jensen's inequality¹ and the fact that $\sum_{i=1}^{x_t} (1/x_t) = 1$, we can give an upper bound for the costs :

style (footnote) okay?

$$x_t * f\left(\frac{\lambda_t}{x_t}\right) \leq x_t \sum_{i=1}^{x_t} \frac{1}{x_t} f(\mu_i \lambda_t) = \frac{x_t}{x_t} \sum_{i=1}^{x_t} f(\mu_i \lambda_t) = \sum_{i=1}^{x_t} f(\mu_i \lambda_t) = \sum_{i \in A} f(\lambda_{i,t})$$

Hence, the operating costs of Σ^* give a lower bound for the operating costs of Σ , and the claim follows. \square

¹For convex $f : \mathbb{R} \rightarrow \mathbb{R}$, arbitrary $\lambda_1, \dots, \lambda_n \in \mathbb{R}$, and $x_1, \dots, x_n \in [0, 1]$ satisfying $\sum_{i=1}^n x_i = 1$ we have:

$$f\left(\sum_{i=1}^n x_i \lambda_i\right) \leq \sum_{i=1}^n x_i f(\lambda_i)$$

Proposition 2.3 tells us that there always exists an optimal schedule that equally distributes its arrival rate to its active servers at any time slot. As a result, we can restrict ourselves in finding such an optimal schedule. Together with corollary 2.2, this allows us to subsequently identify an optimal schedule by its sequence of active servers \mathcal{X} . Moreover, we are now able to simplify our optimisation conditions (1) and (2).

For this, we first define the operating costs function $c_{op}(x, \lambda)$. Given a problem instance \mathcal{I} , $c_{op}(x, \lambda)$ describes the operating costs incurred by equally distributing λ on x active servers using f :

$$c_{op} : \{0, \dots, m\} \times [0, m] \rightarrow \mathbb{R}, \quad c_{op}(x, \lambda) = \begin{cases} 0, & \text{if } x = 0 \\ x * f(\lambda/x), & \text{if } x \neq 0 \wedge \lambda \leq x \\ \infty, & \text{otherwise} \end{cases}$$

Our optimisation conditions now simplify into one single minimalisation:

$$\text{minimise} \quad \underbrace{\sum_{t=1}^T c_{op}(x_t, \lambda_t)}_{\text{operating costs}} + \beta * \underbrace{\sum_{t=1}^T \min\{0, x_t - x_{t-1}\}}_{\text{switching costs}}$$

3 Optimal offline scheduling

In this section, we derive an optimal offline algorithm based on our preliminary work. We reduce our scheduling problem instance \mathcal{I} to a shortest path problem of a level structured graph G . We can then use a dynamic programming approach to find a shortest path of G and subsequently an optimal schedule for \mathcal{I} in pseudo-polynomial time.

3.1 Graph for an optimal schedule

We construct a directed acyclic graph as follows:

$\forall t \in [T - 1]$ and $i, j \in \{0, \dots, m\}$ we add vertices (t, i) modelling the number of active servers at time t . Moreover, we add vertices $(0, 0)$ and $(T, 0)$ for our initial and final state respectively.

In order to warrant that there are at least $\lceil \lambda_t \rceil$ active servers $\forall t \in [T - 1]$, we define an auxiliary function which calculates the costs for handling an arrival rate λ with x active servers:

$$c(x, \lambda) := \begin{cases} 0, & \text{if } x = 0 \\ x f(\lambda/x), & \text{if } x \neq 0 \wedge \lambda \leq x \\ \infty, & \text{otherwise} \end{cases} \quad (4)$$

Then, $\forall t \in [T - 2]$ and $i, j \in \{0, \dots, m\}$, we add edges from (t, i) to $(t + 1, j)$ with weight

$$d(i, j, \lambda_{t+1}) := \beta \max\{0, j - i\} + c(j, \lambda_{t+1}) \quad (5)$$

Finally, for $0 \leq i \leq m$ we add edges from $(0, 0)$ to $(1, i)$ with weight $d(0, i, \lambda_1)$ and from $(T - 1, i)$ to $(T, 0)$ with weight $d(i, 0, \lambda_T) = 0$.

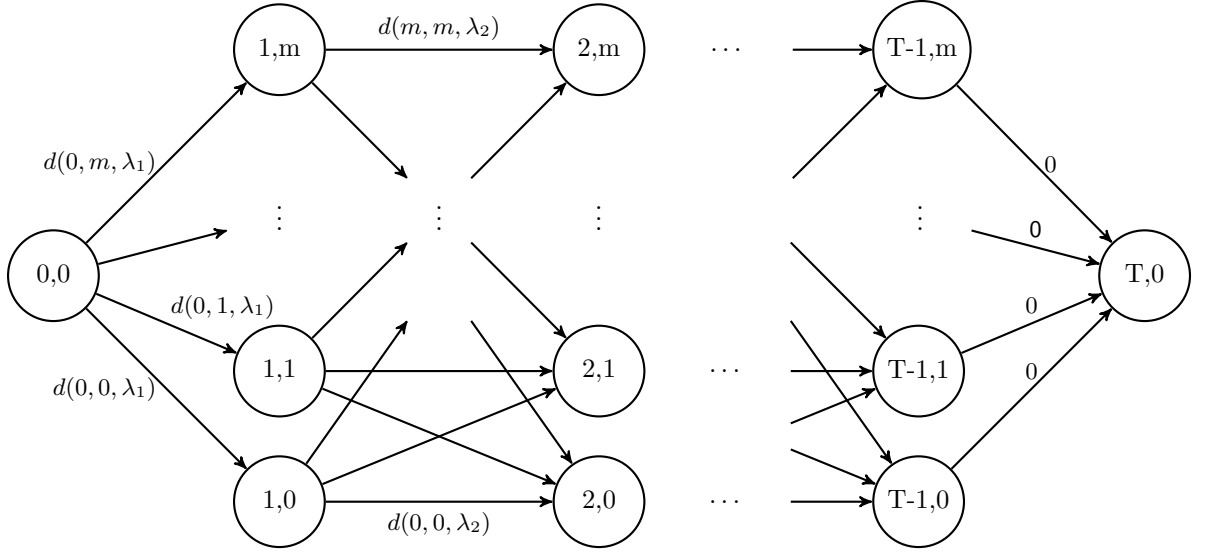


Figure 1: Graph for optimal schedule algorithm.

Note: All edges from (t, i) to $(t + 1, j)$ have weight $d(i, j, \lambda_{t+1})$

Proposition 3.1. *Any given optimal schedule \mathcal{X} corresponds to a shortest path P from $(0, 0)$ to $(T, 0)$ with $\text{costs}(\mathcal{X}) = \text{costs}(P)$ and vice versa.*

Proof.

“ \Rightarrow ”: We construct a feasible path in our graph from \mathcal{X} as follows:

$$\begin{aligned} \text{First set} \quad e_t &:= \left((t, \mathcal{X}(t)), (t + 1, \mathcal{X}(t + 1)) \right), & \forall t \in \{0, \dots, T - 1\} \\ \text{then set} \quad P &:= (e_0, \dots, e_{T-1}) \end{aligned}$$

As each edge e_t in our graph has weight $d(\mathcal{X}(t-1), \mathcal{X}(t), \lambda_t)$, it corresponds to the costs of switching from $\mathcal{X}(t-1)$ to $\mathcal{X}(t)$ servers and processing λ_t with $\mathcal{X}(t)$ active servers. Hence, it directly follows that P is a shortest path of the graph with $\text{costs}(P) = \text{costs}(\mathcal{X})$.

“ \Leftarrow ”: Let $P = ((0, 0) = v_0, \dots, v_T = (T, 0))$ with $v_t \in \{(t, i) \mid 0 \leq i \leq m\}$ be a shortest path of the graph.

We can construct an optimal schedule from P by setting $\mathcal{X} := (v_0(1), \dots, v_T(1))$

By definition (4) it is guaranteed that P only traverses edges such that there are enough active servers $\forall t \in [T]$. Therefore, the created schedule is feasible. Its optimality directly follows from the definition of the edges' weights and so does the equality $\text{costs}(\mathcal{X}) = \text{costs}(P)$.

□

3.2 A pseudo-polynomial minimum cost algorithm

Algorithm 1 Calculate costs for m homogeneous servers

Require: Convex cost function f , $\lambda_0 = \lambda_T = 0$, $\forall t \in [T - 1] : \lambda_t \in [0, m]$

```

1: function SCHEDULE( $m, T, \beta, \lambda_1, \dots, \lambda_{T-1}$ )
2:   if  $T < 2$  then return
3:   let  $p[2 \dots T - 1, m]$  and  $M[1 \dots T - 1, m]$  be new arrays
4:   for  $j \leftarrow 0$  to  $m$  do
5:      $M[1, j] \leftarrow d(0, j, \lambda_1)$ 
6:   for  $t \leftarrow 1$  to  $T - 2$  do
7:     for  $j \leftarrow 0$  to  $m$  do
8:        $opt \leftarrow \infty$ 
9:       for  $i \leftarrow 0$  to  $m$  do
10:         $M[t + 1, j] \leftarrow M[t, i] + d(i, j, \lambda_{t+1})$ 
11:        if  $M[t + 1, j] < opt$  then
12:           $opt \leftarrow M[t + 1, j]$ 
13:           $p[t + 1, j] \leftarrow i$ 
14:         $M[t + 1, j] \leftarrow opt$ 
15:   return  $p$  and  $M$ 

```

Algorithm 2 Extract schedule for m homogeneous servers

```

1: function EXTRACT( $m, p, M, T$ )
2:   let  $x[0 \dots T]$  be a new array
3:    $x[0] \leftarrow x[T] \leftarrow 0$ 
4:   if  $T < 2$  then return  $x$   $\triangleright$  Trivial solution
5:    $x[T - 1] \leftarrow \arg \min_{0 \leq i \leq m} \{M[T - 1, i]\}$ 
6:   for  $t \leftarrow T - 2$  to  $1$  do
7:      $x[t] \leftarrow p[t + 1, x[t + 1]]$ 
8:   return  $x$ 

```

3.2.1 Runtime analysis

The algorithm visits every vertex and every edge of the graph exactly once. As the number of vertices is bounded by $\mathcal{O}(Tm)$ and the number of edges is bounded by $\mathcal{O}(Tm^2)$ the running time is given by:

$$\mathcal{O}(Tm + Tm^2) = \mathcal{O}(Tm^2)$$

As we need $\log_2(m)$ bits to encode m , the running time is polynomial in the numeric value of the input but exponential in the length of the input. Hence, the algorithm is pseudo-polynomial.

3.2.2 A memory optimized algorithm

TODO: use only array with size $2m$

4 A polynomial 4-approximation algorithm for monotonically increasing convex f

We consider a modification of the problem discussed in chapter 3. Assuming that f is convex and monotonically increasing, we can modify our algorithm to obtain a polynomial time 4-approximation algorithm.

4.1 Graph for a 4-optimal schedule

We modify our graph from chapter 3.1 to the reduce the number of vertices. For this, we stop adding m vertices for each timestep, but use vertices that approximate the number of active servers instead. First, let $b := \lceil \log_2(m) \rceil$. We add vertices $(t, 0)$ and $(t, 2^i)$, $\forall t \in [T-1], 0 \leq i \leq b$. All edges and weights are added analogous to chapter 3.1.

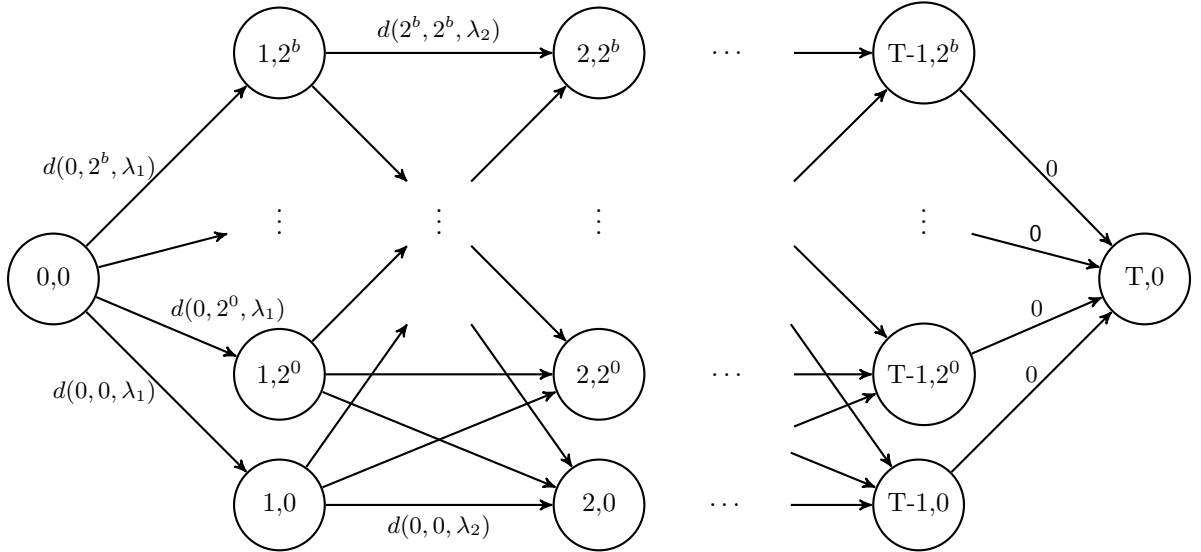


Figure 2: Graph for a 4-approximation algorithm

Definition 4.1. Let $\mathcal{X} = (x_0, \dots, x_T)$ be a schedule and $t > 0$. We say that \mathcal{X} changes its **state** at time t if

$$x_t \neq x_{t-1}$$

and that \mathcal{X} changes its **2-state** at time t if

$$x_t = 0 \quad \text{or} \quad x_t \notin (2^{\lfloor \log_2(x_{t-1}) \rfloor}, 2^{\lceil \log_2(x_{t-1}) \rceil})$$

Proposition 4.2.

1. Any given optimal schedule \mathcal{X} can be transformed to a 4-optimal schedule \mathcal{X}' which corresponds to a path P from $(0, 0)$ to $(T, 0)$ with $\text{costs}(\mathcal{X}') = \text{costs}(P)$.

2. Any shortest path P from $(0, 0)$ to $(T, 0)$ corresponds to a 4-optimal schedule \mathcal{X} with $\text{costs}(P) = \text{costs}(\mathcal{X})$.

Proof.

1. Assume we have an optimal schedule identified by $\mathcal{X} = (x_0, \dots, x_T)$. For $0 \leq t < T$ we inductively set:

$$x'_0 := 0, \quad x'_{t+1} := \begin{cases} \min\{2^{\lceil \log_2(2x_{t+1}) \rceil}, 2^b\}, & \text{if } 0 < x_t \leq x_{t+1} \\ 2^{\lceil \log_2(2x_{t+1}) \rceil}, & \text{if } 0 < x_{t+1} < x_t \text{ and } x'_t \geq 4x_{t+1} \\ x'_t, & \text{if } 0 < x_{t+1} < x_t \text{ and } x'_t < 4x_{t+1} \\ 0, & \text{otherwise} \end{cases} \quad (6)$$

Then let $\mathcal{X}' := (x'_0, \dots, x'_T)$ be the modified sequence of active servers. Notice that $x_t \leq x'_t \leq 4x_t$ holds as x'_t is at most the smallest power of two larger than $2x_t$ which implies that \mathcal{X}' is feasible.

We can now construct a feasible path in our graph from \mathcal{X}' as follows:

$$\begin{aligned} \text{First set} \quad e_t &:= \left((t, \mathcal{X}'(t)), (t+1, \mathcal{X}'(t+1)) \right), \quad \forall t \in \{0, \dots, T-1\} \\ \text{then set} \quad P &:= (e_0, \dots, e_{T-1}) \end{aligned}$$

By the definition of the edges' weights it follows that $\text{costs}(\mathcal{X}') = \text{costs}(P)$.

Next, let $(t_0 = 0, t_1, \dots, t_n = 0)$ be the sequence of times where the optimal schedule \mathcal{X} changes its 2-state. Notice that the modified schedule \mathcal{X}' changes its state only at times t_i and that $2x_{t_i} \leq x'_{t_i}$ holds (TODO: only if not discrete but continuous time steps). This can be seen exemplarily in figure 3 by observing that \mathcal{X}' changes its state only if \mathcal{X} crosses or touches a bordering power of two.

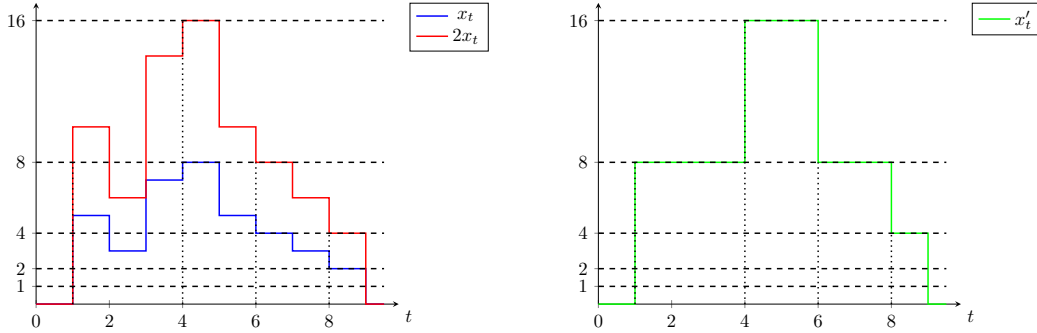


Figure 3: Adaption of an optimal schedule

For this reason, we now only have to consider the fraction of costs of \mathcal{X}' and \mathcal{X} between time steps t_{i-1} and t_i

$$\frac{\text{costs}(\mathcal{X}', t_{i-1}, t_i)}{\text{costs}(\mathcal{X}, t_{i-1}, t_i)} \quad (7)$$

For $x_{t_i} = 0$ it follows from (??) that $\text{costs}(\mathcal{X}', t_{i-1}, t_i) = \text{costs}(\mathcal{X}, t_{i-1}, t_i) = 0$. Hence, we can restrict ourselves to $0 < t_i < T$ with $x_{t_i} \neq 0$. The costs incurred by \mathcal{X}' are given by

$$\begin{aligned}
\text{costs}(\mathcal{X}', t_{i-1}, t_i) &= \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + x'_{t_i} f(\lambda_{t_i}/x'_{t_i}) && \text{by (??)} \\
&\leq \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x'_{t_i}) && \text{by (6)} \\
&\leq \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i}) && \text{f monotonically increasing} \\
\implies \text{costs}(\mathcal{X}', t_{i-1}, t_i) &\leq \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i}) && (8)
\end{aligned}$$

and the costs of \mathcal{X} by

$$\text{costs}(\mathcal{X}, t_{i-1}, t_i) = \beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i}) \quad (9)$$

W.l.o.g. we may assume $x_{t_i} f(\lambda_{t_i}/x_{t_i}) > 0$, otherwise the claim follows trivially. (TODO: is it really trivial?)

(i) $x_{t_i} \leq x_{t_{i-1}}$: From (6) it follows that $x'_{t_i} \leq x'_{t_{i-1}}$. Thus, we can simplify (7):

$$\begin{aligned}
\frac{\text{costs}(\mathcal{X}', t_{i-1}, t_i)}{\text{costs}(\mathcal{X}, t_{i-1}, t_i)} &\leq \frac{\beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by (8),(9)} \\
&= \frac{4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{x_{t_i} f(\lambda_{t_i}/x_{t_i})} && (x_{t_i} \leq x_{t_{i-1}} \text{ and } x'_{t_i} \leq x'_{t_{i-1}}) \\
&= 4
\end{aligned}$$

(ii) $x_{t_i} > x_{t_{i-1}}$: From (6) it follows that $x'_{t_i} \geq x'_{t_{i-1}}$. Thus, we can simplify (7):

$$\begin{aligned}
\frac{\text{costs}(\mathcal{X}', t_{i-1}, t_i)}{\text{costs}(\mathcal{X}, t_{i-1}, t_i)} &\leq \frac{\beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by (8),(9)} \\
&= \frac{\beta(x'_{t_i} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && (x_{t_i} > x_{t_{i-1}} \text{ and } x'_{t_i} \geq x'_{t_{i-1}}) \\
&= \frac{\beta(\min\{2^{\lfloor \log_2(2x_{t_i}) \rfloor}, 2^b\} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by (6)} \\
&\leq \frac{\beta(2^{\lfloor \log_2(2x_{t_i}) \rfloor} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \\
&\leq \frac{\beta(2x_{t_i} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \\
&\leq \frac{\beta(2x_{t_i} - 2x_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by } (2x_{t_{i-1}} \leq x'_{t_{i-1}}) \\
&\leq 4 \frac{\frac{1}{2}\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \\
&\leq 4
\end{aligned}$$

From (i) and (ii) it follows:

$$\text{costs}(\mathcal{X}') \leq 4\text{costs}(\mathcal{X})$$

2. From 1 we obtain that we can construct a 4-optimal path P' from any optimal schedule. Now, let P be a shortest path. We have $costs(P) \leq costs(P') < \infty$, and since every path P with $costs(P) < \infty$ corresponds to a feasible schedule \mathcal{X} with $costs(P) = costs(\mathcal{X})$, \mathcal{X} must also be at least 4-optimal.

□

References

- [1] Minghong Lin, Adam Wierman, Lachlan L. H. Andrew, and Eno Thereska. Dynamic right-sizing for power-proportional data centers. *IEEE/ACM Transactions on Networking (TON)*, 21:1378–1391, 2013.

Appendix

Below, we give an overview of just given definitions and conventions commonly referred to in our paper:

Good idea to have an appendix?

- Input:
 - $m \in \mathbb{N}$... number of homogeneous servers
 - $T \in \mathbb{N}$... number of time slots
 - $\lambda_1, \dots, \lambda_T \in [0, m]$... arrival rates
 - $\Lambda := (\lambda_1, \dots, \lambda_T)$... sequence of arrival rates
 - $\beta \in \mathbb{R}_{\geq 0}$... switching costs of a server
 - $f : [0, 1] \rightarrow \mathbb{R}$... convex operating costs function of a server
 - $\mathcal{I} := (m, T, \Lambda, \beta, f)$... input of a problem instance
- Problem statement:
 - $s_{i,t}$... state of server i at time t , i.e. sleeping (0) or active(1)
 - $S_i := (s_{i,1}, \dots, s_{i,T}) \in \{0, 1\}^T$... sequence of states for server i
 - $\lambda_{i,t} \in [0, 1]$... assigned load for server i at time t
 - $L_i := (\lambda_{i,1}, \dots, \lambda_{i,T}) \in [0, 1]^T$... sequence of assigned loads for server i
 - $\mathcal{S} := (S_1, \dots, S_m)$... sequence of all state changes
 - $\mathcal{L} := (L_1, \dots, L_m)$... sequence of all assigned loads
 - $\Sigma := (\mathcal{S}, \mathcal{L})$... schedule for a problem instance \mathcal{I}
- Miscellaneous:
 - x_t ... number of active servers at time t
 - $\mathcal{X} := (x_1, \dots, x_T)$... sequence of number of active servers
- Conventions:
 - $\lambda_t = 0$ for all $t \notin [T]$, i.e. there is no load before and after the scheduling process
 - $s_{i,t} = 0$ for all $t \notin [T]$, i.e. all servers are powered down before and after the scheduling process