Algorithms for Dynamic Right-Sizing in Data Centers

Kevin Kappelmann

Chair for Theoretical Computer Science,
Technical University of Munich

May 15, 2017

Contents

1 Introduction

TODO: Hardware prices vs. energy costs in data centres, related work and purpose of this paper (offline algorithm, approximation algorithm,...).

1.1 Model description

We want to address the issue of above-mentioned ever-growing energy consumption by examing a scheduling problem that commonly arises in data centres. More specifically, we consider a model consisting of a fixed amount of homogeneous servers denoted by $m \in \mathbb{N}$ and a fixed amount of time slots denoted by $T \in \mathbb{N}$. In turn, each server possesses two power states, i.e. each server is either powered on (*active state*) or powered off (*sleep state*).

For any time slot $t \in [T]$, we have a *mean arrival rate* denoted by λ_t, i.e. the amount of expected load to process in time slot t. We expect the arrival rates to be normalised such that each server can handle a load between 0 and 1 in any time slot. We denote the assigned load for server i in time slot t by $\lambda_{i,t} \in [0,1]$. Consequently, for any time slot t, we expect an arrival rate between 0 and m, that is $\lambda_t \in [0,m]$; otherwise, the servers would not be able to process the given load in time.

The costs incurred by a single machine are described by the sum of the machine's *operating costs*, specified by $f : [0,1] \to \mathbb{R}$, as well as its *(power state) switching costs*, specified by $\beta \in \mathbb{R}_{\geq 0}$. The operating costs f may not exclusively consider energy costs. For example, f may also allow for costs incurred by delays, such as lost revenue caused by users waiting for their responses. Similarly, β may also allow for delay costs, wear and tear costs or the like. [1]

We assume that a sleeping server does not cause any costs. Note that $f(0)$ denotes the costs incurred by an idle server, not a sleeping one; in particular, $f(0)$ may be non-zero. Further, we assume convexity for f. This may seem like a notable restriction at first, but it indeed captures the behaviour of most modern server models. Since we are dealing with homogeneous servers, f and β are the same for all machines.

For convenience, we assume all machines sleeping at time $t = 0$ and force all machines to sleep after the scheduling process, i.e. at times $t > T$. Consequently, every server must power down exactly as many times as it powers on. This allows us to consolidate power up and power down costs into β and to model both costs as being incurred when powering up a server; that is, a model with power up costs $\beta_{\uparrow}$ and power down costs $\beta_{\downarrow}$ can be simply transferred to our model by setting $\beta := \beta'_{\uparrow} := \beta_{\uparrow} + \beta_{\downarrow}$ and $\beta'_{\downarrow} = 0$. Similarly, we assume that there are no loads at times $t \notin [T]$, that is $\lambda_t = \lambda_{i,t} = 0$ for $t \notin [T]$.

1.2 Problem statement

Using above definitions, we can define the input of our model by setting $\mathcal{I} := (m, T, \Lambda, \beta, f)$ where $\Lambda = (\lambda_1, \ldots, \lambda_T)$ is the sequence of arrival rates. We will subsequently identify a problem instance by its input $\mathcal{I}$. Naturally, given a problem instance $\mathcal{I}$, we want to schedule our servers in such a way that we minimise the sum of incurred costs while warranting that we are processing the given loads in time.

For this, consider for each server $i \in [m]$ the sequence of its states S_i and the sequence of its assigned loads L_i; that is

$$S_i := (s_{i,1}, \ldots, s_{i,T}) \in \{0,1\}^T$$
$$L_i := (\lambda_{i,1}, \ldots, \lambda_{i,T}) \in [0,1]^T$$

where $s_{i,t} \in \{0,1\}$ denotes whether server i at time t is sleeping (0) or active (1). Recall that we assume all machines sleeping at times $t \notin [T]$; thus, for $t \notin [T]$ and $i \in [m]$, we have $s_{i,t} = 0$.

We can now define the sequence of all state changes and the sequence of all assigned loads:

$$\mathcal{S} := (S_1, \ldots, S_m)$$
$$\mathcal{L} := (L_1, \ldots, L_m)$$

We will subsequently call a pair $\Sigma := (\mathcal{S}, \mathcal{L})$ a *schedule*. Finally, we are ready to define our problem statement.

Given an input $\mathcal{I}$, our goal is to find a schedule Σ that satisfies the following optimisation:

> Definition of $c(\Sigma)$ too hidden?

$$\text{minimise} \qquad c(\Sigma) := \underbrace{\sum_{t=1}^{T} \sum_{i=1}^{m} \big(f(\lambda_{i,t}) * s_{i,t} \big)}_{\text{operating costs}} + \beta * \underbrace{\sum_{t=1}^{T} \sum_{i=1}^{m} \min\{0, s_{i,t} - s_{i,t-1}\}}_{\text{switching costs}} \qquad (1)$$

$$\text{subject to} \qquad \sum_{i=1}^{m} (\lambda_{i,t} * s_{i,t}) = \lambda_t, \quad \forall t \in [T] \qquad (2)$$

We call a schedule *feasible* if it satisfies (2) and *optimal* if it satisifes (1) and (2).

2 Preliminaries

In this section, we conduct the prepratory work that will lay the foundations for our algorithms. For this, we analyse the structure of feasible schedules concering their cost efficiency in order to find characteristics of optimal schedules; these characteristics will then allow us to greatly simplify our optimisation conditions.

We begin by examing the state sequences of feasible schedules. As we are considering homogeneous servers, we do not care which exact servers process the given work loads. Rather we only care about the amount of active servers and the distribution of loads between them. It is in particular unreasonable to power down a machine and to power on a different machine in return; we could just keep the first machine powered on, saving switching costs. This investigation is captured by our first proposition.

Proposition 2.1. *Given a problem instance $\mathcal{I}$ and a feasible schedule Σ, there exists a feasible schedule Σ' such that*

(i) $c(\Sigma') \leq c(\Sigma)$ and

(ii) Σ' *never powers on and shuts down servers at the same time slot, i.e.* Σ' *satisfies the following formula:*

$$\forall t \in [T]\Big[\big(\forall i \in [m](s_{i,t} - s_{i,t-1} \geq 0)\big) \vee \big(\forall i \in [m](s_{i,t} - s_{i,t-1} \leq 0)\big)\Big] \tag{3}$$

Proof. Let $\Sigma = (\mathcal{S}, \mathcal{L})$ be a feasible schedule for $\mathcal{I}$. We give a procedure that repeatedly modifies Σ such that it satisfies (3) and reduces or retains its costs.

Let $t \in [T]$ be the first time slot falsifying (3). If there does not exist such a time slot, we are done. Otherwise, we can obtain machines $i, j \in [m]$ such that $s_{i,t} - s_{i,t-1} = 1$ and $s_{j,t} - s_{j,t-1} = -1$, i.e. server i powers on at time t and server j powers off. Without loss of generality, we may assume $i < j$.

First, since all servers are sleeping at time $t = 0$, we have

$$s_{k,1} - s_{k,0} = s_{k,1} - 0 = s_{k,1} \geq 0, \quad \forall k \in [m]$$

which satisfies formula (3) for $t = 1$. Thus, we may assume $t > 1$.

Now consider the state sequences of server i and j:

$$S_i = (s_{i,1}, \ldots, s_{i,t-1} = 0, s_{i,t} = 1, \ldots, s_{i,T})$$
$$S_j = (s_{j,1}, \ldots, s_{j,t-1} = 1, s_{j,t} = 0, \ldots, s_{j,T})$$

We modify S_i and S_j by swapping their states for time slots $\geq t$, that is we set

$$S_i' := (s_{i,1}, \ldots, s_{i,t-1} = 0, s_{j,t} = 0, \ldots, s_{j,T})$$
$$S_j' := (s_{j,1}, \ldots, s_{j,t-1} = 1, s_{i,t} = 1, \ldots, s_{i,T})$$

Similarly, we need to swap the assigned loads for server i and j:

$$L_i' := (\lambda_{i,1}, \ldots, \lambda_{i,t-1}, \lambda_{j,t}, \ldots, \lambda_{j,T})$$
$$L_j' := (\lambda_{j,1}, \ldots, \lambda_{j,t-1}, \lambda_{i,t}, \ldots, \lambda_{i,T})$$

Finally, we construct a new schedule $\Sigma' := (\mathcal{S}', \mathcal{L}')$ given by

$$\mathcal{S}' := (S_1, \ldots, S_{i-1}, S_i', S_{i+1}, \ldots, S_{j-1}, S_j', S_{j+1}, \ldots, S_T)$$
$$\mathcal{L}' := (L_1, \ldots, L_{i-1}, L_i', L_{i+1}, \ldots, L_{j-1}, L_j', L_{j+1}, \ldots, L_T)$$

We want to verify that Σ' is a feasible schedule, that is Σ' satisfies (2). For time slots $< t$, the schedules Σ' and Σ still coincide. For time slots $\geq t$, we only changed the order of summation in (2). Thus, Σ' is feasible.

Σ and Σ' coincide in their operating costs; however, their switching costs differ in that there are no switching costs β at time t for server i using Σ'. As we assume $\beta \geq 0$, we conclude $c(\Sigma') \leq c(\Sigma)$

Moreover, we decreased the amount of bad spots at time t concerning (3). Hence, by re- peating described process on Σ', we obtain a terminating procedure that returns a schedule satisfying the conditions. $\qquad \square$

bad spots? better description?

3

As a special case, proposition 2.1 unfolds its power on optimal schedules, yielding the next corrolary.

Corollary 2.2. *Given a problem instance $\mathcal{I}$, there exists an optimal schedule Σ^* satisfying* (3).

Proof. Let Σ be an optimal schedule for $\mathcal{I}$. Applying proposition 2.1 to Σ yields Σ^* and the claim follows. $\qquad\square$

Next, we want to consider the sequence of active servers. For this, let $\mathcal{X}$ denote the sequence of sums of active servers at each time slot t, that is

$$\mathcal{X} := \left(x_1 = \sum_{i=1}^{m} s_{i,1}, \ldots, x_T = \sum_{i=1}^{m} s_{i,T}\right) \in \{0, \ldots, m\}^T$$

As we assume all machines sleeping at times $t \notin [T]$, we have $x_t = 0$ for $t \notin [T]$.

The next proposition poses the cornerstone of our subsequent works. We want to establish an optimal scheduling strategy given a fixed amount of active servers. It turns out that equal load-sharing seems a very desirable strategy.

Proposition 2.3 (Equal load sharing). *Given $x_t \in \mathbb{N}$ active servers in time slot t, an arrival rate $\lambda_t \in [0, x_t]$, and a convex cost function f, a most cost-efficient and feasible scheduling strategy is to assign each active server a load of λ_t/x_t.*

Proof. Let Σ be an arbitrary, feasible schedule using x_t servers in time slot t, and let A be its set of active servers in time slot t, that is $A := \{i \in [m] \mid s_{i,t} = 1\}$. Consider the operating costs of Σ at time t given by

$$\sum_{i=1}^{m} \big(f(\lambda_{i,t}) * s_{i,t}\big) = \sum_{i \in A} \big(f(\lambda_{i,t}) * 1\big) + \sum_{i \in [m] \setminus A} \big(f(\lambda_{i,t}) * 0\big) = \sum_{i \in A} f(\lambda_{i,t})$$

Since Σ is feasible (see constraint (2)), we have

$$\sum_{i \in A} \lambda_{i,t} = \lambda_t$$

Hence, we can obtain weights $\mu_1, \ldots, \mu_{x_t} \in [0, 1]$ that relate $\lambda_{i,t}$ and λ_t for $i \in A$ such that

$$\sum_{i=1}^{x_t} \mu_i = 1 \quad \text{and} \quad \sum_{i \in A} f(\lambda_{i,t}) = \sum_{i=1}^{x_t} f(\mu_i \lambda_t) \tag{4}$$

In particular, we have

$$\sum_{i=1}^{x_t} \mu_i \lambda_t = \lambda_t \tag{5}$$

Using these weights, we now consider the operating costs of a schedule Σ^* that equally distributes λ_t to its x_t active servers:

$$\sum_{i=1}^{x_t} f\left(\frac{\lambda_t}{x_t}\right) = x_t * f\left(\frac{\lambda_t}{x_t}\right) \overset{(5)}{=} x_t * f\left(\sum_{i=1}^{x_t} \frac{\mu_i \lambda_t}{x_t}\right)$$

4

With the use of Jensen's inequality[1] and the fact that $\sum_{i=1}^{x_t}(1/x_t) = 1$, we can give an upper bound for the costs :

$$x_t * f\left(\frac{\lambda_t}{x_t}\right) \le x_t \sum_{i=1}^{x_t} \frac{1}{x_t} f(\mu_i \lambda_t) = \frac{x_t}{x_t} \sum_{i=1}^{x_t} f(\mu_i \lambda_t) = \sum_{i=1}^{x_t} f(\mu_i \lambda_t) \stackrel{(4)}{=} \sum_{i \in A} f(\lambda_{i,t})$$

Thus, the operating costs of Σ^* give a lower bound for the operatings costs of Σ, and the claim follows. $\qquad\square$

As a special case, we can again apply our just derived proposition to optimal schedules.

Corollary 2.4. *Given a problem instance $\mathcal{I}$, there exists an optimal schedule Σ^* that equally distributes its arrival rates to its active servers in each time slot.*

Proof. Let $\Sigma = (\mathcal{S}, \mathcal{L})$ be an optimal schedule for $\mathcal{I}$. We exchange $\mathcal{L}$ with a new strategy $\mathcal{L}^*$ that equally distributes the arrival rates to all active servers of Σ in each time slot. Setting $\Sigma^* := (\mathcal{S}, \mathcal{L}^*)$ we have $c(\Sigma^*) \le c(\Sigma)$ by proposition 2.3 and the claim follows. $\qquad\square$

As a result of corrolary 2.4, we can restrict ourselves in finding an optimal schedule that equally distributes its arrival rates to its active servers. Together with corollary 2.2, this allows us to subsequently identify an optimal schedule by its sequence of active servers $\mathcal{X}$.

We are now able to simplify our optimisation conditions (1) and (2). For this, given a problem instance $\mathcal{I}$, we define the operating costs function $c_{op}(x, \lambda)$ that describes the costs incurred by equally distributing λ on x active servers using f:

$$c_{op} : \{0, \ldots, m\} \times [0, m] \to \mathbb{R} \cup \{\infty\}, \quad c_{op}(x, \lambda) = \begin{cases} 0, & \text{if } x = 0 \\ x * f(\lambda/x), & \text{if } x \ne 0 \land \lambda \le x \\ \infty, & \text{if } x \ne 0 \land \lambda > x \end{cases}$$

We assign infinite costs in case $\lambda > x$ as there would be too few active servers to process the arrival rate, i.e. the schedule would not be feasible. Next, we define the switching costs function $c_{sw}(x_{t-1}, x_t)$ that describes the incurred costs when changing the amount of active server from x_{t-1} to x_t:

$$c_{sw}(x_{t-1}, x_t) := \beta * \max\{0, x_t - x_{t-1}\}$$

Lastly, we can define the costs function $c(x_{t-1}, x_t, \lambda_t)$ that describes the incurring costs for a single time step using an equal distribution of loads:

$$c(x_{t-1}, x_t, \lambda_t) := c_{op}(x_t, \lambda_t) + c_{sw}(x_{t-1}, x_t)$$

---

[1] For convex $f : \mathbb{R} \to \mathbb{R}$, arbitrary $\lambda_1, \ldots, \lambda_n \in \mathbb{R}$, and $x_1, \ldots, x_n \in [0, 1]$ satisfying $\sum_{i=1}^{n} x_i = 1$ we have:

$$f\left(\sum_{i=1}^{n} x_i \lambda_i\right) \le \sum_{i=1}^{n} x_i f(\lambda_i)$$

The optimisation conditions for a schedule now simplify to one single minimalisation:

$$\text{minimise} \quad c(\mathcal{X}) := \sum_{t=1}^{T} c(x_{t-1}, x_t, \lambda_t) \tag{6}$$

We subsequently call a schedule $\mathcal{X}$ *optimal* if it satisfies (6).

3   Optimal offline scheduling

In this section, we derive two optimal offline algorithms based on our preliminary work. First, we reduce our problem specified by $\mathcal{I}$ to a shortest path problem of a level structured graph G. We then use a dynamic programming approach to find a shortest path in G and thereby an optimal schedule for $\mathcal{I}$ in pseudo-polynomial time $\Theta(Tm^2)$. After that, we refine our initial approach to derive an improved algorithm with pseudo-linear complexity $\Theta(Tm)$.

> Sounds strange...

3.1   A pseudo-polynomial algorithm

Let $\mathcal{I}$ be a problem instance. Thanks to our preliminary work, we know that there exists an optimal schedule which is identifiable by its sequence of active servers $\mathcal{X}$. In order to find this sequence $\mathcal{X}$, consider the weighted, level structured graph G defined as follows:

$$
\begin{aligned}
V &:= \left\{ v_{x,t} \mid x \in \{0, \dots, m\}, t \in [T] \right\} \uplus \{v_{0,0}, v_{0,T+1}\} \\
E &:= \left\{ (v_{x,t}, v_{x',t+1}) \mid x, x' \in \{0, \dots, m\}, t \in \{0, \dots, T\}, v_{x,t}, v_{x',t+1} \in V \right\} \\
c_G(v_{x,t}, v_{x',t+1}) &:= c(x, x', \lambda_{t+1}) \\
G &:= (V, E, c_G)
\end{aligned}
$$

For any possible amount of active servers x and any time slot t, we add a node $v_{x,t}$. Moreover, we add a start node $v_{0,0}$ as well as an end node $v_{0,T+1}$. Next, we connect all nodes to their successors with respect to time. Semantically, $v_{x,t}$ denotes the state of scheduling the arrival rate λ_t equally to x servers in time slot t. For any edge connecting $v_{x,t}$ with $v_{x',t+1}$, we assign costs $c(x, x', \lambda_{t+1})$, i.e. the costs of taking the edge correspond to switching from x to x' machines and processing the load λ_{t+1} with x' machines.

> Maybe something different than c_G?

The costs of a path $P = (v_{x_0,T_0}, v_{x_1,T_0+1}, \dots, v_{x_n,T_0+n})$ with length n in our graph (where $x_t \in \{0, \dots, m\}, T_0 \in \{0, \dots, T\}$ and $T_0 + n \leq T + 1$) are thus given by

> Is this well-written?

$$c(P) := \sum_{t=1}^{n} c(x_{t-1}, x_t, \lambda_{T_0+t})$$

In particular, for a path $P = (v_{0,0}, v_{x_1,1}, \dots, v_{x_T,T}, v_{0,T+1})$ from our start node $v_{0,0}$ to our end node $v_{0,T+1}$ we have

$$c(P) = c(0, x_1, \lambda_1) + \sum_{t=2}^{T} c(x_{t-1}, x_t, \lambda_t) + \underbrace{c(x_T, 0, 0)}_{=0} = c(0, x_1, \lambda_1) + \sum_{t=2}^{T} c(x_{t-1}, x_t, \lambda_t) \tag{7}$$

6

Note that the costs of such a path directly correspond to those of a schedule $\mathcal{X}$ (see (6)). Any shortest path from $v_{0,0}$ to $v_{0,T+1}$ is thus forced to minimise the costs of the corresponding schedule. Needless to say, this demands for a proof of correctness.


Figure 1: Level structured graph for a pseudo-polynomial, optimal offline algorithm

Lemma 3.1. *Any given schedule $\mathcal{X}$ for $\mathcal{I}$ corresponds to a path P from $v_{0,0}$ to $v_{0,T+1}$ with $c(\mathcal{X}) = c(P)$ and vice versa.*

Proof.

"$\Rightarrow$": Let $\mathcal{X} = (x_1, \ldots, x_T)$ be a schedule for $\mathcal{I}$. We construct a path in our graph by setting

$$P := (v_{0,0}, v_{x_1,1}, v_{x_2,2}, \ldots, v_{x_T,T}, v_{0,T+1})$$

We examine the costs and conclude

$$c(\mathcal{X}) \overset{(6)}{=} \sum_{t=1}^{T} c(x_{t-1}, x_t, \lambda_t) = \underbrace{c(x_0, x_1, \lambda_1)}_{=c(0, x_1, \lambda_1)} + \sum_{t=2}^{T} c(x_{t-1}, x_t, \lambda_t) \overset{(7)}{=} c(P)$$

"$\Leftarrow$": Let $P = (v_{0,0}, v_{x_1,1}, v_{x_2,2}, \ldots, v_{x_T,T}, v_{0,T+1})$ be a path from $v_{0,0}$ to $v_{0,T+1}$. We construct a schedule for $\mathcal{I}$ by setting

$$\mathcal{X} := (x_1, \ldots, x_T)$$

Again, by examing the costs we conclude

$$c(P) \overset{(7)}{=} \underbrace{c(0, x_1, \lambda_1)}_{=c(x_0, x_1, \lambda_1)} + \sum_{t=2}^{T} c(x_{t-1}, x_t, \lambda_t) = \sum_{t=1}^{T} c(x_{t-1}, x_t, \lambda_t) \overset{(6)}{=} c(\mathcal{X})$$

Thus, the claim follows. □

Theorem 3.2. *Any given optimal schedule $\mathcal{X}$ for $\mathcal{I}$ corresponds to a shortest path P from $v_{0,0}$ to $v_{0,T+1}$ with $c(\mathcal{X}) = c(P)$ and vice versa.*

Proof. By lemma 3.1 we have a one-to-one correspondence between schedules $\mathcal{X}$ and paths P obeying $c(\mathcal{X}) = c(P)$. Thus, we have

$$c(\mathcal{X}) \text{ minimal} \iff c(P) \text{ minimal}$$

and the claim follows. □

In the following, we give an algorithm based on our just verified construction. We split our procedure into two subroutines.

SHORTEST_PATHS uses a dynamic programming approach similar to the well-known Bellman-Ford algorithm. It returns the minimum costs to all nodes as well as the predecessor of any node in respect to its shortest path.

EXTRACT_SCHEDULE uses the predecessors calculated by SHORTEST_PATHS in order to obtain the sequence of nodes describing a shortest path and thereby a schedule for our problem instance.

Algorithm 1 Pseudo-polynomial optimal offline scheduling

1: **function** OPTIMAL_OFFLINE_SCHEDULING$(m, T, \Lambda, \beta, f)$
2:   $(C, P) \leftarrow$ SHORTEST_PATHS$(m, T, \Lambda, \beta, f)$
3:   $\mathcal{X} \leftarrow$ EXTRACT_SCHEDULE(P, T)
4:   **return** $\mathcal{X}$

5: **function** SHORTEST_PATHS$(m, T, \Lambda, \beta, f)$
6:   **let** $C[T+1, m]$ and $P[T+1, m]$ **be** new arrays   ▷ Costs and predecessors of nodes
7:   **for** $x \leftarrow 0$ to m **do**   ▷ Initialisation
8:    $C[1, x] \leftarrow c(0, x, \lambda_1)$
9:    $P[1, x] \leftarrow 0$
10:   **for** $t \leftarrow 2$ to T **do**   ▷ Iterative calculation of costs and predecessors
11:    **for** $x' \leftarrow 0$ to m **do**
12:     $P[t, x'] \leftarrow \underset{x \in \{0,\ldots,m\}}{\arg\min} \left\{ C[t-1, x] + c(x, x', \lambda_t) \right\}$ ▷ Find best preceding choice
13:     $C[t, x'] \leftarrow C\big[t-1, P[t, x']\big] + c\big(P[t, x'], x', \lambda_t\big)$
14:   $P[T+1, 0] \leftarrow \underset{x \in \{0,\ldots,m\}}{\arg\min} \left\{ C[T, x] \right\}$ ▷ Find best choice for last time slot
15:   $C[T+1, 0] \leftarrow C\big[T, P[T+1, 0]\big]$
16:   **return** (C, P)

17: **function** EXTRACT_SCHEDULE(P, T)
18:   **let** $\mathcal{X}[T]$ **be** a new array
19:   $\mathcal{X}[T] \leftarrow P[T+1, 0]$ ▷ Get best choice for last time slot
20:   **for** $t \leftarrow T-1$ to 1 **do**   ▷ Iteratively obtain a schedule by using the predecessors
21:    $\mathcal{X}[t] \leftarrow P\big[t+1, \mathcal{X}[t+1]\big]$
22:   **return** $\mathcal{X}$

Naturally, we are interested in our algorithm's time and memory complexity. Subroutine SHORTEST_PATHS needs $\Theta(m)$ iterations for its initialisation, $\Theta(Tm^2)$ steps for the iterative calculation, and $\Theta(m)$ steps for its final minimisation search. In addition, EXTRACT_SCHEDULE needs $\Theta(T)$ iterations for its schedule retrieval. Thus, we receive a runtime of

$$\Theta(m + Tm^2 + m + T) = \Theta(Tm^2)$$

The runtime is polynomial in the numeric value of the input; however, as we just need $\log_2(m)$ bits to encode m, it is exponential in the length of the input. Hence, the algorithm is pseudo-polynomial.

Our memory demand is defined by the size of the arrays C and P. As both arrays are of size $\Theta(Tm)$, we have a memory complexity of $\Theta(Tm)$

3.2   A pseudo-linear algorithm

The algorithm developed in section 3.1 is of a quite simple nature. Its underlying graph G is able to represent any possible schedule $\mathcal{X}$ since we simply we add an edge for any possible scheduling choice at any possible time slot. This approach seems rather intuitive and readily verifiable, but this convenience comes with a cost; the density of G causes a quadratic runtime in the number of servers m. In order to improve the runtime to pseudo-linear complexity, we need to thin out our graph.

Let us revise our initial algorithm. Our graph consists of nodes $v_{x,t}$. Any node $v_{x,t}$ denotes the state of scheduling the arrival rate λ_t equally to x servers in time slot t. The algorithm calculates the minimum costs to all nodes. The cost of a node $v_{x,t}$ thus corresponds to the lowest achievable cost up to time slot t of all schedules $\mathcal{X}$ that assign x servers at time t to process the arrival rate λ_t. In particular, the cost of the end node $v_{0,T+1}$ tell us the minimum cost of all schedules. This approach, however, does not consider the possibility to schedule $y \neq x$ servers in time slot t and to switch to x servers just at the very last moment of t when calculating the cost of $v_{x,t}$. Consider the following two examples:

(a) State $v_{0,1}$ could be reached with cost 3 by moving down from node $v_{1,1}$.

(b) State $v_{1,1}$ could be reached with cost 1 by moving up from $v_{0,1}$.

Figure 2: Two examples depicting a shortcoming of our initial approach. The calculated nodes' costs are highlighted in red. Dashed edges are not part of the initial algorithm.

9

Although our algorithm delivers the correct end results, its immediate steps are somewhat unsatisfying. We want our states to capture a more general notion than given in section 3.1; preferably, we would like a node $v_{x,t}$ to denote the state of having x active servers at the end of time slot t. In practice, we may reach a state $v_{x,t}$ by moving down from a state $v_{y^\downarrow,t}$ where $y^\downarrow > i$ with cost 0 or by moving up from a state $v_{y^\uparrow,t}$ where $y^\uparrow < i$ with cost $\beta * (x - y^\uparrow)$.

In order to allow for these new possibilities, we define a directed, weighted graph as follows:

$$V := \{v_{x,t\downarrow} \mid x \in \{0, \ldots, m\}, t \in [T]\} \cup \{v_{x,t\uparrow} \mid x \in \{0, \ldots, m\}, t \in [T]\} \cup \{v_{0,0}, v_{0,T+1}\}$$
$$E_s := \{(v_{0,0}, v_{x,1\downarrow}) \mid x \in \{0, \ldots, m\}\}$$
$$E_e := \{(v_{x,T\uparrow}, v_{0,T+1}) \mid x \in \{0, \ldots, m\}\}$$
$$E_\downarrow := \{(v_{x,t\downarrow}, v_{x-1,t\downarrow}) \mid x \in [m], t \in [T]\}$$
$$E_\uparrow := \{(v_{x-1,t\uparrow}, v_{x,t\uparrow}) \mid x \in [m], t \in [T]\}$$
$$E_{\downarrow\uparrow} := \{(v_{x,t\downarrow}, v_{x,t\uparrow}) \mid x \in \{0, \ldots, m\}, t \in [T]\}$$
$$E_{\uparrow\downarrow} := \{(v_{x,t\uparrow}, v_{x,t+1\downarrow}) \mid x \in \{0, \ldots, m\}, t \in [T]\}$$
$$E := E_s \cup E_e \cup E_\downarrow \cup E_\uparrow \cup E_{\downarrow\uparrow} \cup E_{\uparrow\downarrow}$$
$$c_G(e) := \begin{cases} c(0, x, \lambda_1), & \text{if } e = (v_{0,0}, v_{x,1}) \in E_s \\ c_{op}(x, \lambda_{t+1}), & \text{if } e = (v_{x,t}, v_{x,t+1}) \in E_{\uparrow\downarrow} \\ \beta, & \text{if } e \in E_\uparrow \\ 0, & \text{if } e \in (E_\downarrow \cup E_{\downarrow\uparrow} \cup E_e) \end{cases}$$
$$G := (V, E, c_G)$$

As the phrase goes, *"A picture is worth a thousand words."* See figure 3 for a more appealing illustration.

For any possible amount of active servers x and any time slot t, we add two nodes $v_{x,t\downarrow}$ and $v_{x,t\uparrow}$. Semantically, the cost of $v_{x,t\downarrow}$ will denote the minimum cost up to time slot t when processing λ_t with x or more servers whereas the cost of $v_{x,t\uparrow}$ will denote the minimum cost of having x active servers at the end of time slot t. Moreover, we again add a start node $v_{0,0}$ as well as an end node $v_{0,T+1}$.

The set of edges E_s and E_e denote the start initialisation and the end minimisation step, respectively. An edge $(v_{x,t\uparrow}, v_{x,t+1\downarrow}) \in E_{\uparrow\downarrow}$ accounts for the operating costs that incur when processing the arrival rate λ_{t+1} with x active servers.

After every time step from $t-1$ to t that incurs operatings costs, we have a minimisation step in our graph. For this, we first move down along the chain $v_{m,t\downarrow}, v_{m-1,t\downarrow}, \ldots, v_{0,t\downarrow}$ using edges from $E_\downarrow$ with cost 0. Then we proceed to move to the right from $v_{0,t\downarrow}$ to $v_{0,t\uparrow}$. Lastly, we move up along the chain $v_{0,t\uparrow}, v_{1,t\uparrow}, \ldots, v_{m,t\uparrow}$ using edges from $E_\uparrow$ with cost β. In order to have the possibility to keep the calculated costs of $v_{x,t\downarrow}$ during the upward movement, we added edges $(v_{x,t\downarrow}, v_{x,t\uparrow}) \in E_{\downarrow\uparrow}$ with cost 0.

This minimisation step is the key to our runtime improvement. It facilitates the determination of the best predecessor of a state $v_{x,t\downarrow}$ because we already know that the minimum cost of having x servers at the end of time slot $t-1$ is stored in $v_{x,t-1\uparrow}$. Thus, the cheapest possibility to process the next arrival rate λ_t using x servers can simply be calculated by adding $c_{op}(x, \lambda_t)$ to the cost of $v_{x,t-1\uparrow}$. Consequently, the cost of $v_{x,t\downarrow}$ is given by the minimum of $v_{x,t-1\uparrow} + c_{op}(x, \lambda_t)$ and $v_{x+1,\downarrow}$.
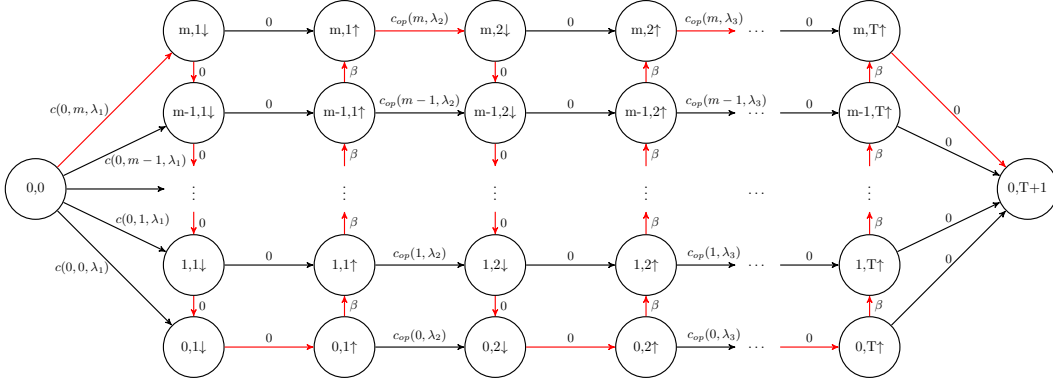
caption of figure

Figure 3: Graph for a pseudo-linear, optimal offline algorithm; the path of the topological sorting is highlighted in red.

As one can see in above figure, we stretched our graph but at the same time also greatly reduced the amount of edges compared to our initial approach in section 3.1. By following the marked path of the topological sorting, we can work our way through the graph to calculate the shortest paths, ultimately reaching the destination $v_{0,T+1}$. Our next task shall be the verification of our new construction.

Proposition 3.3. *Any given schedule $\mathcal{X}$ for $\mathcal{I}$ corresponds to a path P from $v_{0,0}$ to $v_{0,T+1}$ with $c(\mathcal{X}) = c(P)$.*

Proof. Let $\mathcal{X} = (x_1, \ldots, x_T)$ be a schedule for $\mathcal{I}$. $\qquad\qquad\qquad\qquad\qquad\qquad\square$

4   A polynomial 4-approximation algorithm for monotonically increasing convex f

We consider a modification of the problem discussed in chapter. Assuming that f is convex and monotonically increasing, we can modify our algorithm to obtain a polynomial time 4-approximation algorithm.

4.1   Graph for a 4-optimal schedule

We modify our graph from section 3.1 to the reduce the number of nodes. For this, we stop adding m nodes for each timestep, but use nodes that approximate the number of active servers instead. First, let $b := \lceil \log_2(m) \rceil$. We add nodes $(t,0)$ and $(t,2^i), \forall t \in [T-1], 0 \le i \le b$. All edges and weights are added analogous to chapter 3.1.

Figure 4: Graph for a 4-approximation algorithm

Definition 4.1. Let $\mathcal{X} = (x_0, \ldots, x_T)$ be a schedule and $t > 0$.
We say that $\mathcal{X}$ changes its **state** at time t if

$$x_t \neq x_{t-1}$$

and that $\mathcal{X}$ changes its **2-state** at time t if

$$x_t = 0 \quad \text{or} \quad x_t \notin \left(2^{\lfloor \log_2(x_{t-1}) \rfloor}, 2^{\lceil \log_2(x_{t-1}) \rceil} \right)$$

Proposition 4.2.

1. *Any given optimal schedule $\mathcal{X}$ can be transformed to a 4-optimal schedule $\mathcal{X}'$ which corresponds to a path P from $(0,0)$ to $(T,0)$ with $costs(\mathcal{X}') = costs(P)$.*

2. *Any shortest path P from $(0,0)$ to $(T,0)$ corresponds to a 4-optimal schedule $\mathcal{X}$ with $costs(P) = costs(\mathcal{X})$.*

Proof.

1. Assume we have an optimal schedule identified by $\mathcal{X} = (x_0, \ldots, x_T)$. For $0 \leq t < T$ we inductively set:

$$x_0' := 0, \qquad x_{t+1}' := \begin{cases} \min\left\{ 2^{\lfloor \log_2(2x_{t+1}) \rfloor}, 2^b \right\}, & \text{if } 0 < x_t \leq x_{t+1} \\ 2^{\lceil \log_2(2x_{t+1}) \rceil}, & \text{if } 0 < x_{t+1} < x_t \text{ and } x_t' \geq 4x_{t+1} \\ x_t', & \text{if } 0 < x_{t+1} < x_t \text{ and } x_t' < 4x_{t+1} \\ 0, & \text{otherwise} \end{cases} \tag{8}$$

Then let $\mathcal{X}' := (x_0', \ldots, x_T')$ be the modified sequence of active servers. Notice that $x_t \leq x_t' \leq 4x_t$ holds as x_t' is at most the smallest power of two larger than $2x_t$ which

12

implies that $\mathcal{X}'$ is feasible.

We can now construct a feasible path in our graph from $\mathcal{X}'$ as follows:

First set $e_t := \Big((t, \mathcal{X}'(t)), (t+1, \mathcal{X}'(t+1))\Big), \qquad \forall t \in \{0, \dots, T-1\}$

then set $P := (e_0, \dots, e_{T-1})$

By the definition of the edges' weights it follows that $costs(\mathcal{X}') = costs(P)$.

Next, let $(t_0 = 0, t_1, \dots, t_n = 0)$ be the sequence of times where the optimal schedule $\mathcal{X}$ changes its 2-state. Notice that the modified schedule $\mathcal{X}'$ changes its state only at times t_i and that $2x_{t_i} \le x'_{t_i}$ holds (TODO: only if not discrete but continous time steps). This can be seen exemplarily in figure 5 by obvserving that $\mathcal{X}'$ changes its state only if $\mathcal{X}$ crosses or touches a bordering power of two.


Figure 5: Adaption of an optimal schedule

For this reason, we now only have to consider the fraction of costs of $\mathcal{X}'$ and $\mathcal{X}$ between time steps t_{i-1} and t_i

$$\frac{costs(\mathcal{X}', t_{i-1}, t_i)}{costs(\mathcal{X}, t_{i-1}, t_i)} \tag{9}$$

For $x_{t_i} = 0$ it follows from (??) that $costs(\mathcal{X}', t_{i-1}, t_i) = costs(\mathcal{X}, t_{i-1}, t_i) = 0$. Hence, we can restrict ourselves to $0 < t_i < T$ with $x_{t_i} \ne 0$. The costs incurred by $\mathcal{X}'$ are given by

$$costs(\mathcal{X}', t_{i-1}, t_i) = \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + x'_{t_i} f(\lambda_{t_i}/x'_{t_i}) \qquad \text{by (??)}$$

$$\le \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x'_{t_i}) \qquad \text{by (8)}$$

$$\le \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i}) \quad \text{f monotonically increasing}$$

$$\implies \quad costs(\mathcal{X}', t_{i-1}, t_i) \le \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i}) \tag{10}$$

and the costs of $\mathcal{X}$ by

$$costs(\mathcal{X}, t_{i-1}, t_i) = \beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i}) \tag{11}$$

W.l.o.g. we may assume $x_{t_i} f(\lambda_{t_i}/x_{t_i}) > 0$, otherwise the claim follows trivially. (TODO: is it really trivial?)

(i) $\underline{x_{t_i} \leq x_{t_{i-1}}}$: From (8) it follows that $x'_{t_i} \leq x'_{t_{i-1}}$. Thus, we can simplify (9):

$$\frac{costs(\mathcal{X}', t_{i-1}, t_i)}{costs(\mathcal{X}, t_{i-1}, t_i)} \leq \frac{\beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \qquad \text{by (10),(11)}$$

$$= \frac{4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{x_{t_i} f(\lambda_{t_i}/x_{t_i})} \qquad (x_{t_i} \leq x_{t_{i-1}} \text{ and } x'_{t_i} \leq x'_{t_{i-1}})$$

$$= 4$$

(ii) $\underline{x_{t_i} > x_{t_{i-1}}}$: From (8) it follows that $x'_{t_i} \geq x'_{t_{i-1}}$. Thus, we can simplify (9):

$$\frac{costs(\mathcal{X}', t_{i-1}, t_i)}{costs(\mathcal{X}, t_{i-1}, t_i)} \leq \frac{\beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \qquad \text{by (10),(11)}$$

$$= \frac{\beta(x'_{t_i} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \qquad (x_{t_i} > x_{t_{i-1}} \text{ and } x'_{t_i} \geq x'_{t_{i-1}})$$

$$= \frac{\beta(\min\{2^{\lfloor \log_2(2x_{t_i})\rfloor}, 2^b\} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \qquad \text{by (8)}$$

$$\leq \frac{\beta(2^{\lfloor \log_2(2x_{t_i})\rfloor} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})}$$

$$\leq \frac{\beta(2x_{t_i} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})}$$

$$\leq \frac{\beta(2x_{t_i} - 2x_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \qquad \text{by } (2x_{t_{i-1}} \leq x'_{t_{i-1}})$$

$$\leq 4 \frac{\frac{1}{2}\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})}$$

$$\leq 4$$

From (i) and (ii) it follows:

$$costs(\mathcal{X}') \leq 4costs(\mathcal{X})$$

2. From 1 we obtain that we can construct a 4-optimal path P' from any optimal schedule. Now, let P be a shortest path. We have $costs(P) \leq costs(P') < \infty$, and since every path P with $costs(P) < \infty$ corresponds to a feasible schedule $\mathcal{X}$ with $costs(P) = costs(\mathcal{X})$, $\mathcal{X}$ must also be at least 4-optimal.

$\square$

References

[1] Minghong Lin, Adam Wierman, Lachlan L. H. Andrew, and Eno Thereska. Dynamic right-sizing for power-propotional data centers. *IEEE/ACM Transactions on Networking (TON)*, 21:1378–1391, 2013.

Appendix

Below, we give an overview of just given definitions and conventions commonly referred to in our paper:

> Good idea to have an appendix?

- Input:
    - $m \in \mathbb{N}$ ... number of homogeneous servers
    - $T \in \mathbb{N}$ ... number of time slots
    - $\lambda_1, \ldots, \lambda_T \in [0, m]$ ... arrival rates
    - $\Lambda \coloneqq (\lambda_1, \ldots, \lambda_T)$ ... sequence of arrival rates
    - $\beta \in \mathbb{R}_{\geq 0}$ ... switching costs of a server
    - $f : [0, 1] \to \mathbb{R}$ ... convex operating costs function of a server
    - $\mathcal{I} \coloneqq (m, T, \Lambda, \beta, f)$ ... input of a problem instance

- Problem statement:
    - $s_{i,t} \in \{0, 1\}$ ... state of server i at time t, i.e. sleeping (0) or active(1)
    - $S_i \coloneqq (s_{i,1}, \ldots, s_{i,T})$ ... sequence of states for server i
    - $\lambda_{i,t} \in [0, 1]$ ... assigned load for server i at time t
    - $L_i \coloneqq (\lambda_{i,1}, \ldots, \lambda_{i,T})$ ... sequence of assigned loads for server i
    - $\mathcal{S} \coloneqq (S_1, \ldots, S_m)$ ... sequence of all state changes
    - $\mathcal{L} \coloneqq (L_1, \ldots, L_m)$ ... sequence of all assigned loads
    - $\Sigma \coloneqq (\mathcal{S}, \mathcal{L})$ ... schedule for a problem instance $\mathcal{I}$

- Miscellaneous:
    - $x_t \in \{0, \ldots, m\}$ ... number of active servers at time t
    - $\mathcal{X} \coloneqq (x_1, \ldots, x_T)$ ... sequence of number of active servers

- Conventions:
    - $\lambda_t = 0$ for all $t \notin [T]$, i.e. there is no load before and after the scheduling process
    - $s_{i,t} = 0$ for all $t \notin [T]$, i.e. all servers are powered down before and after the scheduling process