

# Algorithms for Dynamic Right-Sizing in Data Centers

Kevin Kappelmann

*Chair for Theoretical Computer Science,  
Technical University of Munich*

May 8, 2017

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Model description . . . . .	1
1.2	Problem statement . . . . .	1
<b>2</b>	<b>Preliminaries</b>	<b>2</b>
<b>3</b>	<b>Optimal scheduling for m homogeneous servers</b>	<b>5</b>
3.1	Graph for an optimal schedule . . . . .	5
3.2	A pseudo-polynomial minimum cost algorithm . . . . .	7
3.2.1	Runtime analysis . . . . .	7
3.2.2	A memory optimized algorithm . . . . .	7
<b>4</b>	<b>A polynomial 4-approximation algorithm for monotonically increasing convex f</b>	<b>8</b>
4.1	Graph for a 4-optimal schedule . . . . .	8

# 1 Introduction

TODO: Hardware prices vs. energy costs in data centres and purpose of this paper (offline algorithm, approximation algorithm,...).

## 1.1 Model description

We want to address the issue of the above-mentioned ever-growing energy consumption by examining a scheduling problem commonly arising in data centres. More specifically, we consider a model consisting of a fixed amount of homogeneous servers denoted by  $m \in \mathbb{N}$  and a fixed amount of time slots denoted by  $T \in \mathbb{N}$ . In turn, each server possesses two power states, i.e. each server is either powered on (*active state*) or powered off (*sleep state*).

Better name than sleep state?

For any time slot  $t \in [T]$  we have a *mean arrival rate* denoted by  $\lambda_t$ , i.e. the amount of expected load to process in time slot  $t$ . We expect the arrival rates to be normalised such that each server  $i \in [m]$  can handle a load between zero and one in any time slot. We denote the assigned load for server  $i$  in time slot  $t$  by  $\lambda_{i,t} \in [0, 1]$ . Consequently, for any time slot  $t$  we expect an arrival rate between 0 and  $m$ , i.e.  $\lambda_t \in [0, m]$ ; otherwise, we would not be able to process the given load in time.

The incurred energy costs of a single machine can be described by the sum of the machine's (*power state*) *switching costs* specified by  $\beta \in \mathbb{R}_{\geq 0}$  as well as its *operating costs* specified by  $f : [0, 1] \rightarrow \mathbb{R}$ . We assume that a sleeping server does not generate any energy costs. Note that  $f(0)$  describes the costs generated by an idle server, not a sleeping one; in particular,  $f(0)$  may be unequal to zero. Further, we assume convexity for  $f$ . This may seem like a notable restriction at first, but it indeed captures the behaviour of most modern server models. Since we are dealing with homogeneous servers,  $\beta$  and  $f$  are the same for all machines.

We want to stress that  $f$  may not only pose as a depiction of energy costs. For example,  $f$  may also allow for costs incurred by delays, such as lost revenue caused by users that need to wait for their responses. Similarly,  $\beta$  may also allow for delay costs, wear and tear costs or the like. [1]

For convenience, we assume all machines sleeping at time  $t = 0$  and force all machines to sleep after the scheduling process, i.e. at times  $t > T$ . This justifies the consolidation of power up and power down costs into  $\beta$  because it allows us to model both costs as being incurred when powering up a server; that is, a model with power up costs  $\beta_{\uparrow}$  and power down costs  $\beta_{\downarrow}$  can simply be transferred to our model by setting  $\beta := \beta_{\uparrow} + \beta_{\downarrow}$ . We can now proceed to define our problem statement.

## 1.2 Problem statement

Using above definitions, we can define the input of our model by  $\mathcal{I} := (m, T, \Lambda, \beta, f)$  where  $\Lambda = (\lambda_1, \dots, \lambda_T)$  is the sequence of arrival rates. Naturally, given an input  $\mathcal{I}$ , we want to schedule our servers in such a way that we minimise the sum of incurred costs while warranting that we are processing the given loads in time.

For this, we consider for each server  $i \in [m]$  the sequence of its states  $S_i$  and the sequence of its assigned loads  $L_i$ , that is

$$S_i := (s_{i,1}, \dots, s_{i,T}) \in \{0, 1\}^T$$

$$L_i := (\lambda_{i,1}, \dots, \lambda_{i,T}) \in [0, 1]^T$$

where  $s_{i,t} \in \{0, 1\}$  denotes whether server  $i$  at time  $t$  is sleeping (0) or active (1). Recall that we assume all machines sleeping at times  $t \notin [T]$ ; consequently, for  $t \notin [T]$  and  $i \in [m]$  we have  $s_{i,t} = 0$ .

Using these sequences  $S_i$  and  $L_i$ , we can now define the sequence of all state changes and the sequence of all assigned loads:

good sentence?  
 $\in (\{0, 1\}^m)^T$

$$\mathcal{S} := (S_1, \dots, S_m)$$

$$\mathcal{L} := (L_1, \dots, L_m)$$

We will subsequently call a pair  $\Sigma := (\mathcal{S}, \mathcal{L})$  a *schedule*. Given an input  $\mathcal{I}$ , our goal is to find a schedule  $\Sigma$  that satisfies the following optimisation:

Definition of  $c(\Sigma)$   
too hidden?

$$\text{minimise} \quad c(\Sigma) := \underbrace{\sum_{t=1}^T \sum_{i=1}^m (f(\lambda_{i,t}) * s_{i,t})}_{\text{operating costs}} + \beta * \underbrace{\sum_{t=1}^T \sum_{i=1}^m \min\{0, s_{i,t} - s_{i,t-1}\}}_{\text{switching costs}} \quad (1)$$

$$\text{subject to} \quad \sum_{i=1}^m (\lambda_{i,t} * s_{i,t}) = \lambda_t, \quad \forall t \in [T] \quad (2)$$

We call a schedule *feasible* if it satisfies (2) and *optimal* if it satisfies (1) and (2).

## 2 Preliminaries

In this section, we conduct the preparatory work needed for our algorithms and proofs.

This sounds strange... What should be written here?

**Proposition 2.1.** *Given a problem instance  $\mathcal{I}$  and a feasible schedule  $\Sigma$ , there exists a feasible schedule  $\Sigma'$  such that*

(i)  $c(\Sigma') \leq c(\Sigma)$  and

(ii)  $\Sigma'$  never powers on and shuts down servers at the same time slot, i.e.  $\Sigma'$  satisfies the following formula:

$$\forall t \in [T] \left[ (\forall i \in [m] (s_{i,t} - s_{i,t-1} \geq 0)) \vee (\forall i \in [m] (s_{i,t} - s_{i,t-1} \leq 0)) \right] \quad (3)$$

*Proof.* Let  $\Sigma = (\mathcal{S}, \mathcal{L})$  be a feasible schedule for  $\mathcal{I}$ . We give a procedure that repeatedly modifies  $\Sigma$  such that it satisfies (3) and reduces or retains its costs.

Let  $t \in [T]$  be the first time slot falsifying (3). If there does not exist such a  $t$ , we are done. Otherwise, we can obtain  $i, j \in [m]$  such that  $s_{i,t} - s_{i,t-1} = 1$  and  $s_{j,t} - s_{j,t-1} = -1$ . WLOG we may assume  $i < j$ .

First, since all servers are sleeping at time  $t = 0$ , we have

$$s_{k,1} - s_{k,0} = s_{k,1} - 0 = s_{k,1} \geq 0, \quad \forall k \in [m]$$

Thus, we may assume  $t > 1$ . Now consider the state sequences  $S_i$  and  $S_j$ :

$$\begin{aligned} S_i &= (s_{i,1}, \dots, s_{i,t-1} = 0, s_{i,t} = 1, \dots, s_{i,T}) \\ S_j &= (s_{j,1}, \dots, s_{j,t-1} = 1, s_{j,t} = 0, \dots, s_{j,T}) \end{aligned}$$

We modify  $S_i$  and  $S_j$  by swapping their states for time slots  $\geq t$ , that is we set

$$\begin{aligned} S'_i &:= (s_{i,1}, \dots, s_{i,t-1} = 0, s_{j,t} = 0, \dots, s_{j,T}) \\ S'_j &:= (s_{j,1}, \dots, s_{j,t-1} = 1, s_{i,t} = 1, \dots, s_{i,T}) \end{aligned}$$

Similarly, we need to swap the assigned loads:

$$\begin{aligned} L'_i &:= (\lambda_{i,1}, \dots, \lambda_{i,t-1}, \lambda_{j,t}, \dots, \lambda_{j,T}) \\ L'_j &:= (\lambda_{j,1}, \dots, \lambda_{j,t-1}, \lambda_{i,t}, \dots, \lambda_{i,T}) \end{aligned}$$

Finally, we construct a new schedule  $\Sigma' := (\mathcal{S}', \mathcal{L}')$  given by

$$\begin{aligned} \mathcal{S}' &:= (S_1, \dots, S'_i, \dots, S'_j, \dots, S_T) \\ \mathcal{L}' &:= (L_1, \dots, L'_i, \dots, L'_j, \dots, L_T) \end{aligned}$$

We now want to verify that  $\Sigma'$  is a feasible schedule, that is  $\Sigma'$  satisfies (2). For time slots  $< t$  the schedules  $\Sigma'$  and  $\Sigma$  still coincide. For time slots  $\geq t$  we only changed the order of summation (2). Thus,  $\Sigma'$  is feasible.

Further,  $\Sigma'$  and  $\Sigma$  coincide in their operating costs; however, their switching costs differ in that there are no switching costs at time slot  $t$  for server  $i$  using  $\Sigma'$ . As we assume  $\beta \geq 0$ , we conclude  $c(\Sigma') \leq c(\Sigma)$ .

Moreover, we decreased the amount of bad spots at time slot  $t$  concerning (3). Hence, by repeating described process on  $\Sigma'$ , we obtain a terminating procedure that returns a schedule satisfying the conditions.  $\square$

bad spots? better description?

**Corollary 2.2.** *Given a problem instance  $\mathcal{I}$ , there exists an optimal schedule  $\Sigma$  satisfying (3).*

*Proof.* Let  $\Sigma$  be an optimal schedule for  $\mathcal{I}$ . Applying proposition 2.1 to  $\Sigma$  yields the result.  $\square$

Next, we want to consider the sequence of active servers. For this, let  $\mathcal{X}$  denote the sequence of sums of active servers at each time slot  $t$ , i.e.

$$\mathcal{X} := (x_1 = \sum_{i=1}^m s_{i,1}, \dots, x_T = \sum_{i=1}^m s_{i,T}) \in \{0, \dots, m\}^T$$

Recall that we assume all machines sleeping at times  $t \notin [T]$ ; consequently, for  $t \notin [T]$  we have  $x_t = 0$ .

**Proposition 2.3** (Equal load sharing). *Given a schedule  $\Sigma$  with  $x_t \in \mathbb{N}$  active servers at time slot  $t$ , an arrival rate  $\lambda_t \in [0, x_t]$ , and a convex cost function  $f$ , a most cost-efficient and feasible scheduling strategy is to assign each active server a load of  $\lambda_t/x_t$ .*

This is sounds awkwardly formulated, doesn't it?

*Proof.* Let  $A$  be the set of active servers at time slot  $t$ , that is  $A := \{i \in [m] \mid s_{i,t} = 1\}$  (note that  $|A| = x_t$ ). Consider the schedule's operating costs at time slot  $t$  given by

$$\sum_{i=1}^m (f(\lambda_{i,t}) * s_{i,t}) = \sum_{i \in A} f(\lambda_{i,t})$$

Since we are only interested in feasible schedules (see constraint (2)), we have

$$\sum_{i \in A} \lambda_{i,t} = \lambda_t$$

Thus, we can obtain weights  $\mu_1, \dots, \mu_{x_t} \in [0, 1]$  such that

$$\begin{aligned} \sum_{i=1}^{x_t} \mu_i &= 1 \\ \text{and} \quad \sum_{i \in A} f(\lambda_{i,t}) &= \sum_{i=1}^{x_t} f(\mu_i \lambda_t) \end{aligned}$$

Using these weights, we now consider the operating costs of a schedule that equally distributes  $\lambda_t$  on its  $x_t$  active servers:

$$\begin{aligned} \sum_{i=1}^{x_t} f\left(\frac{\lambda_t}{x_t}\right) &= \sum_{i=1}^{x_t} f\left(\sum_{i=1}^{x_t} \frac{\mu_i \lambda_t}{x_t}\right) \\ \implies \sum_{i=1}^{x_t} f\left(\frac{\lambda_t}{x_t}\right) &\leq \sum_{i=1}^{x_t} \sum_{i=1}^{x_t} \frac{1}{x_t} f(\mu_i \lambda_t) \quad (\text{by Jensen's inequality}) \\ \iff x_t * f\left(\frac{\lambda_t}{x_t}\right) &\leq \sum_{i=1}^{x_t} \frac{1}{x_t} \sum_{i=1}^{x_t} f(\mu_i \lambda_t) \\ \iff x_t * f\left(\frac{\lambda_t}{x_t}\right) &\leq \sum_{i=1}^{x_t} f(\mu_i \lambda_t) \end{aligned}$$

Hence, equally distributing our load gives us a lower bound for our costs, and the claim follows.  $\square$

Proposition 2.3 tells us that there always exists an optimal schedule that equally shares its arrival rate to its active servers at any time slot. As a result, we can restrict ourselves in finding such an optimal schedule. Together with corollary 2.2, this allows us to subsequently identify an optimal schedule by its sequence of active servers  $\mathcal{X}$ . Moreover, we are now able to simplify our optimisation conditions (1) and (2).

For this, we first define the operating costs function  $c_{op}(\cdot, \cdot)$  of a problem instance that gives us the costs of scheduling a load  $\lambda$  equally on  $x$  active servers using  $f$ :

$$c_{op} : \{0, \dots, m\} \times [0, m] \rightarrow \mathbb{R}, \quad (x, \lambda) \mapsto \begin{cases} 0, & \text{if } x = 0 \\ x * f(\lambda/x), & \text{if } x \neq 0 \wedge \lambda \leq x \\ \infty, & \text{otherwise} \end{cases}$$

Our optimisation conditions now simplify into one single minimalisation:

$$\text{minimise} \quad \underbrace{\sum_{t=1}^T c_{op}(x_t, \lambda_t)}_{\text{operating costs}} + \beta * \underbrace{\sum_{t=1}^T \min\{0, x_t - x_{t-1}\}}_{\text{switching costs}}$$

### 3 Optimal scheduling for m homogeneous servers

TODO: introduction text

#### 3.1 Graph for an optimal schedule

We construct a directed acyclic graph as follows:

$\forall t \in [T - 1]$  and  $i, j \in \{0, \dots, m\}$  we add vertices  $(t, i)$  modelling the number of active servers at time  $t$ . Moreover, we add vertices  $(0, 0)$  and  $(T, 0)$  for our initial and final state respectively.

In order to warrant that there are at least  $\lceil \lambda_t \rceil$  active servers  $\forall t \in [T - 1]$ , we define an auxiliary function which calculates the costs for handling an arrival rate  $\lambda$  with  $x$  active servers:

$$c(x, \lambda) := \begin{cases} 0, & \text{if } x = 0 \\ x f(\lambda/x), & \text{if } x \neq 0 \wedge \lambda \leq x \\ \infty, & \text{otherwise} \end{cases} \quad (4)$$

Then,  $\forall t \in [T - 2]$  and  $i, j \in \{0, \dots, m\}$ , we add edges from  $(t, i)$  to  $(t + 1, j)$  with weight

$$d(i, j, \lambda_{t+1}) := \beta \max\{0, j - i\} + c(j, \lambda_{t+1}) \quad (5)$$

Finally, for  $0 \leq i \leq m$  we add edges from  $(0, 0)$  to  $(1, i)$  with weight  $d(0, i, \lambda_1)$  and from  $(T - 1, i)$  to  $(T, 0)$  with weight  $d(i, 0, \lambda_T) = 0$ .

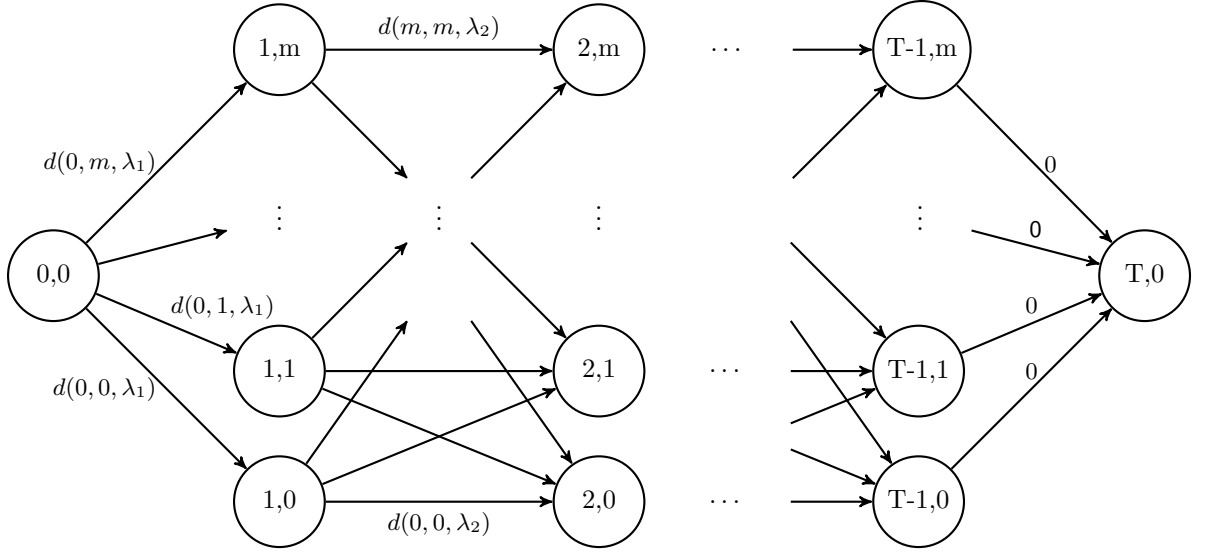


Figure 1: Graph for optimal schedule algorithm.

**Note:** All edges from  $(t, i)$  to  $(t + 1, j)$  have weight  $d(i, j, \lambda_{t+1})$

**Proposition 3.1.** *Any given optimal schedule  $\mathcal{X}$  corresponds to a shortest path  $P$  from  $(0, 0)$  to  $(T, 0)$  with  $\text{costs}(\mathcal{X}) = \text{costs}(P)$  and vice versa.*

*Proof.*

“ $\Rightarrow$ ”: We construct a feasible path in our graph from  $\mathcal{X}$  as follows:

$$\begin{aligned} \text{First set} \quad e_t &:= \left( (t, \mathcal{X}(t)), (t + 1, \mathcal{X}(t + 1)) \right), & \forall t \in \{0, \dots, T - 1\} \\ \text{then set} \quad P &:= (e_0, \dots, e_{T-1}) \end{aligned}$$

As each edge  $e_t$  in our graph has weight  $d(\mathcal{X}(t-1), \mathcal{X}(t), \lambda_t)$ , it corresponds to the costs of switching from  $\mathcal{X}(t-1)$  to  $\mathcal{X}(t)$  servers and processing  $\lambda_t$  with  $\mathcal{X}(t)$  active servers. Hence, it directly follows that  $P$  is a shortest path of the graph with  $\text{costs}(P) = \text{costs}(\mathcal{X})$ .

“ $\Leftarrow$ ”: Let  $P = ((0, 0) = v_0, \dots, v_T = (T, 0))$  with  $v_t \in \{(t, i) \mid 0 \leq i \leq m\}$  be a shortest path of the graph.

We can construct an optimal schedule from  $P$  by setting  $\mathcal{X} := (v_0(1), \dots, v_T(1))$

By definition (4) it is guaranteed that  $P$  only traverses edges such that there are enough active servers  $\forall t \in [T]$ . Therefore, the created schedule is feasible. Its optimality directly follows from the definition of the edges' weights and so does the equality  $\text{costs}(\mathcal{X}) = \text{costs}(P)$ .

□



### 3.2 A pseudo-polynomial minimum cost algorithm

---

**Algorithm 1** Calculate costs for  $m$  homogeneous servers

---

**Require:** Convex cost function  $f$ ,  $\lambda_0 = \lambda_T = 0$ ,  $\forall t \in [T - 1] : \lambda_t \in [0, m]$

```

1: function SCHEDULE( $m, T, \beta, \lambda_1, \dots, \lambda_{T-1}$ )
2:   if  $T < 2$  then return
3:   let  $p[2 \dots T - 1, m]$  and  $M[1 \dots T - 1, m]$  be new arrays
4:   for  $j \leftarrow 0$  to  $m$  do
5:      $M[1, j] \leftarrow d(0, j, \lambda_1)$ 
6:   for  $t \leftarrow 1$  to  $T - 2$  do
7:     for  $j \leftarrow 0$  to  $m$  do
8:        $opt \leftarrow \infty$ 
9:       for  $i \leftarrow 0$  to  $m$  do
10:         $M[t + 1, j] \leftarrow M[t, i] + d(i, j, \lambda_{t+1})$ 
11:        if  $M[t + 1, j] < opt$  then
12:           $opt \leftarrow M[t + 1, j]$ 
13:           $p[t + 1, j] \leftarrow i$ 
14:         $M[t + 1, j] \leftarrow opt$ 
15:   return  $p$  and  $M$ 

```

---



---

**Algorithm 2** Extract schedule for  $m$  homogeneous servers

---

```

1: function EXTRACT( $m, p, M, T$ )
2:   let  $x[0 \dots T]$  be a new array
3:    $x[0] \leftarrow x[T] \leftarrow 0$ 
4:   if  $T < 2$  then return  $x$  ▷ Trivial solution
5:    $x[T - 1] \leftarrow \arg \min_{0 \leq i \leq m} \{M[T - 1, i]\}$ 
6:   for  $t \leftarrow T - 2$  to  $1$  do
7:      $x[t] \leftarrow p[t + 1, x[t + 1]]$ 
8:   return  $x$ 

```

---

#### 3.2.1 Runtime analysis

The algorithm visits every vertex and every edge of the graph exactly once. As the number of vertices is bounded by  $\mathcal{O}(Tm)$  and the number of edges is bounded by  $\mathcal{O}(Tm^2)$  the running time is given by:

$$\mathcal{O}(Tm + Tm^2) = \mathcal{O}(Tm^2)$$

As we need  $\log_2(m)$  bits to encode  $m$ , the running time is polynomial in the numeric value of the input but exponential in the length of the input. Hence, the algorithm is pseudo-polynomial.

#### 3.2.2 A memory optimized algorithm

TODO: use only array with size  $2m$

## 4 A polynomial 4-approximation algorithm for monotonically increasing convex $f$

We consider a modification of the problem discussed in chapter 3. Assuming that  $f$  is convex and monotonically increasing, we can modify our algorithm to obtain a polynomial time 4-approximation algorithm.

### 4.1 Graph for a 4-optimal schedule

We modify our graph from chapter 3.1 to the reduce the number of vertices. For this, we stop adding  $m$  vertices for each timestep, but use vertices that approximate the number of active servers instead. First, let  $b := \lceil \log_2(m) \rceil$ . We add vertices  $(t, 0)$  and  $(t, 2^i)$ ,  $\forall t \in [T-1], 0 \leq i \leq b$ . All edges and weights are added analogous to chapter 3.1.




Figure 2: Graph for a 4-approximation algorithm

**Definition 4.1.** Let  $\mathcal{X} = (x_0, \dots, x_T)$  be a schedule and  $t > 0$ . We say that  $\mathcal{X}$  changes its **state** at time  $t$  if

$$x_t \neq x_{t-1}$$

and that  $\mathcal{X}$  changes its **2-state** at time  $t$  if

$$x_t = 0 \quad \text{or} \quad x_t \notin (2^{\lfloor \log_2(x_{t-1}) \rfloor}, 2^{\lceil \log_2(x_{t-1}) \rceil})$$

**Proposition 4.2.**

1. Any given optimal schedule  $\mathcal{X}$  can be transformed to a 4-optimal schedule  $\mathcal{X}'$  which corresponds to a path  $P$  from  $(0, 0)$  to  $(T, 0)$  with  $\text{costs}(\mathcal{X}') = \text{costs}(P)$ .

2. Any shortest path  $P$  from  $(0, 0)$  to  $(T, 0)$  corresponds to a 4-optimal schedule  $\mathcal{X}$  with  $\text{costs}(P) = \text{costs}(\mathcal{X})$ .

*Proof.*

1. Assume we have an optimal schedule identified by  $\mathcal{X} = (x_0, \dots, x_T)$ . For  $0 \leq t < T$  we inductively set:

$$x'_0 := 0, \quad x'_{t+1} := \begin{cases} \min\{2^{\lceil \log_2(2x_{t+1}) \rceil}, 2^b\}, & \text{if } 0 < x_t \leq x_{t+1} \\ 2^{\lceil \log_2(2x_{t+1}) \rceil}, & \text{if } 0 < x_{t+1} < x_t \text{ and } x'_t \geq 4x_{t+1} \\ x'_t, & \text{if } 0 < x_{t+1} < x_t \text{ and } x'_t < 4x_{t+1} \\ 0, & \text{otherwise} \end{cases} \quad (6)$$

Then let  $\mathcal{X}' := (x'_0, \dots, x'_T)$  be the modified sequence of active servers. Notice that  $x_t \leq x'_t \leq 4x_t$  holds as  $x'_t$  is at most the smallest power of two larger than  $2x_t$  which implies that  $\mathcal{X}'$  is feasible.

We can now construct a feasible path in our graph from  $\mathcal{X}'$  as follows:

$$\begin{aligned} \text{First set} \quad e_t &:= \left( (t, \mathcal{X}'(t)), (t+1, \mathcal{X}'(t+1)) \right), \quad \forall t \in \{0, \dots, T-1\} \\ \text{then set} \quad P &:= (e_0, \dots, e_{T-1}) \end{aligned}$$

By the definition of the edges' weights it follows that  $\text{costs}(\mathcal{X}') = \text{costs}(P)$ .

Next, let  $(t_0 = 0, t_1, \dots, t_n = 0)$  be the sequence of times where the optimal schedule  $\mathcal{X}$  changes its 2-state. Notice that the modified schedule  $\mathcal{X}'$  changes its state only at times  $t_i$  and that  $2x_{t_i} \leq x'_{t_i}$  holds (TODO: only if not discrete but continuous time steps). This can be seen exemplarily in figure 3 by observing that  $\mathcal{X}'$  changes its state only if  $\mathcal{X}$  crosses or touches a bordering power of two.

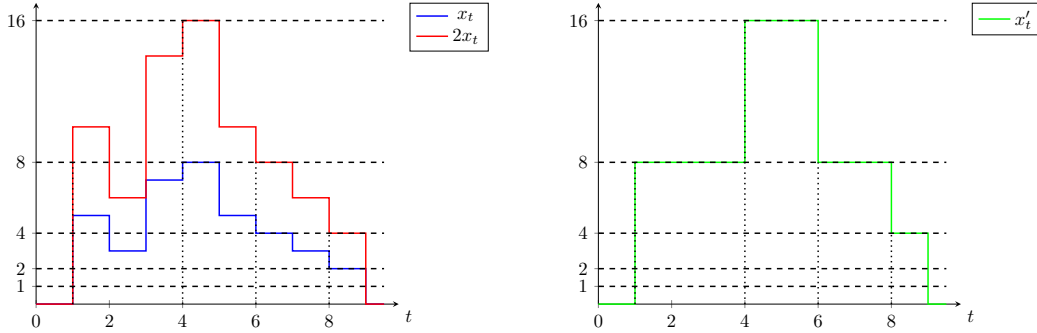


Figure 3: Adaption of an optimal schedule

For this reason, we now only have to consider the fraction of costs of  $\mathcal{X}'$  and  $\mathcal{X}$  between time steps  $t_{i-1}$  and  $t_i$

$$\frac{\text{costs}(\mathcal{X}', t_{i-1}, t_i)}{\text{costs}(\mathcal{X}, t_{i-1}, t_i)} \quad (7)$$

For  $x_{t_i} = 0$  it follows from (??) that  $\text{costs}(\mathcal{X}', t_{i-1}, t_i) = \text{costs}(\mathcal{X}, t_{i-1}, t_i) = 0$ . Hence, we can restrict ourselves to  $0 < t_i < T$  with  $x_{t_i} \neq 0$ . The costs incurred by  $\mathcal{X}'$  are given by

$$\begin{aligned}
\text{costs}(\mathcal{X}', t_{i-1}, t_i) &= \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + x'_{t_i} f(\lambda_{t_i}/x'_{t_i}) && \text{by (??)} \\
&\leq \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x'_{t_i}) && \text{by (6)} \\
&\leq \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i}) && \text{f monotonically increasing} \\
\implies \text{costs}(\mathcal{X}', t_{i-1}, t_i) &\leq \beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i}) && (8)
\end{aligned}$$

and the costs of  $\mathcal{X}$  by

$$\text{costs}(\mathcal{X}, t_{i-1}, t_i) = \beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i}) \quad (9)$$

W.l.o.g. we may assume  $x_{t_i} f(\lambda_{t_i}/x_{t_i}) > 0$ , otherwise the claim follows trivially. (TODO: is it really trivial?)

(i)  $x_{t_i} \leq x_{t_{i-1}}$ : From (6) it follows that  $x'_{t_i} \leq x'_{t_{i-1}}$ . Thus, we can simplify (7):

$$\begin{aligned}
\frac{\text{costs}(\mathcal{X}', t_{i-1}, t_i)}{\text{costs}(\mathcal{X}, t_{i-1}, t_i)} &\leq \frac{\beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by (8),(9)} \\
&= \frac{4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{x_{t_i} f(\lambda_{t_i}/x_{t_i})} && (x_{t_i} \leq x_{t_{i-1}} \text{ and } x'_{t_i} \leq x'_{t_{i-1}}) \\
&= 4
\end{aligned}$$

(ii)  $x_{t_i} > x_{t_{i-1}}$ : From (6) it follows that  $x'_{t_i} \geq x'_{t_{i-1}}$ . Thus, we can simplify (7):

$$\begin{aligned}
\frac{\text{costs}(\mathcal{X}', t_{i-1}, t_i)}{\text{costs}(\mathcal{X}, t_{i-1}, t_i)} &\leq \frac{\beta \max\{0, x'_{t_i} - x'_{t_{i-1}}\} + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta \max\{0, x_{t_i} - x_{t_{i-1}}\} + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by (8),(9)} \\
&= \frac{\beta(x'_{t_i} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && (x_{t_i} > x_{t_{i-1}} \text{ and } x'_{t_i} \geq x'_{t_{i-1}}) \\
&= \frac{\beta(\min\{2^{\lfloor \log_2(2x_{t_i}) \rfloor}, 2^b\} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by (6)} \\
&\leq \frac{\beta(2^{\lfloor \log_2(2x_{t_i}) \rfloor} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \\
&\leq \frac{\beta(2x_{t_i} - x'_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \\
&\leq \frac{\beta(2x_{t_i} - 2x_{t_{i-1}}) + 4x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} && \text{by } (2x_{t_{i-1}} \leq x'_{t_{i-1}}) \\
&\leq 4 \frac{\frac{1}{2}\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})}{\beta(x_{t_i} - x_{t_{i-1}}) + x_{t_i} f(\lambda_{t_i}/x_{t_i})} \\
&\leq 4
\end{aligned}$$

From (i) and (ii) it follows:

$$\text{costs}(\mathcal{X}') \leq 4\text{costs}(\mathcal{X})$$

2. From 1 we obtain that we can construct a 4-optimal path  $P'$  from any optimal schedule. Now, let  $P$  be a shortest path. We have  $costs(P) \leq costs(P') < \infty$ , and since every path  $P$  with  $costs(P) < \infty$  corresponds to a feasible schedule  $\mathcal{X}$  with  $costs(P) = costs(\mathcal{X})$ ,  $\mathcal{X}$  must also be at least 4-optimal.

□

## References

- [1] Minghong Lin, Adam Wierman, Lachlan L. H. Andrew, and Eno Thereska. Dynamic right-sizing for power-proportional data centers. *IEEE/ACM Transactions on Networking (TON)*, 21:1378–1391, 2013.

## Appendix

Below, we give an overview of just given definitions and conventions commonly referred to in our paper:

Good idea to have an appendix?

- Input:
  - $m \in \mathbb{N}$ ... number of homogeneous servers
  - $T \in \mathbb{N}$ ... number of time slots
  - $\lambda_1, \dots, \lambda_T \in [0, m]$ ... arrival rates
  - $\Lambda := (\lambda_1, \dots, \lambda_T)$ ... sequence of arrival rates
  - $\beta \in \mathbb{R}_{\geq 0}$ ... switching costs of a server
  - $f : [0, 1] \rightarrow \mathbb{R}$ ... convex operating costs function of a server
  - $\mathcal{I} := (m, T, \Lambda, \beta, f)$ ... input of a problem instance
- Problem statement:
  - $s_{i,t}$ ... state of server  $i$  at time  $t$ , i.e. sleeping (0) or active(1)
  - $S_i := (s_{i,1}, \dots, s_{i,T}) \in \{0, 1\}^T$ ... sequence of states for server  $i$
  - $\lambda_{i,t} \in [0, 1]$ ... assigned load for server  $i$  at time  $t$
  - $L_i := (\lambda_{i,1}, \dots, \lambda_{i,T}) \in [0, 1]^T$ ... sequence of assigned loads for server  $i$
  - $\mathcal{S} := (S_1, \dots, S_m)$ ... sequence of all state changes
  - $\mathcal{L} := (L_1, \dots, L_m)$ ... sequence of all assigned loads
  - $\Sigma := (\mathcal{S}, \mathcal{L})$ ... schedule for a problem instance  $\mathcal{I}$
- Miscellaneous:
  - $x_t$ ... number of active servers at time  $t$
  - $\mathcal{X} := (x_1, \dots, x_T)$ ... sequence of number of active servers
- Conventions:
  - $\lambda_t = 0$  for all  $t \notin [T]$ , i.e. there is no load before and after the scheduling process
  - $s_{i,t} = 0$  for all  $t \notin [T]$ , i.e. all servers are powered down before and after the scheduling process